

Data-Driven Discovery of Vehicle Dynamics Using Probabilistic Bayesian Neural Networks

Chinmay Samak

csamak@clemson.edu

Tanmay Samak

tsamak@clemson.edu

Vasanth Seethapathi

vasants@clemson.edu



Responsibility Assignment

- Data Collection and Preprocessing:
 - Tanmay
 - Vasanth
- Problem Formulation:
 - Chinmay
 - Tanmay
- Implementation:
 - Tanmay
 - Chinmay
- Hyperparameter Tuning and Pipeline Optimization:
 - Everyone

Note: Responsibility does not indicate contribution. We have no conflicts of interest to declare.



Motivation & Objective

- **Motivation:**

- Real-time multi-step dynamic simulation of complex multi-body models
- Deep learning - learn implicit dynamics and generalize across multiple operating conditions
- Bayesian inference - quantify uncertainty in model predictions for explainability & reliability

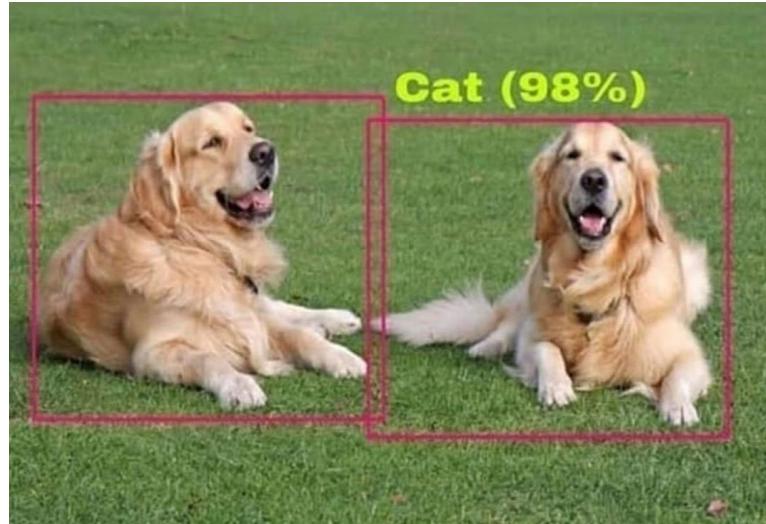
- **Objectives:**

- Forward simulation for state propagation
- Replace hardware sensor by software model
- State-estimation in coverage denied areas or in sensor malfunction conditions
- Redundant variable measurement for more robust state-estimation
- Predicting anomalies in sensor measurements and fault-tolerant control

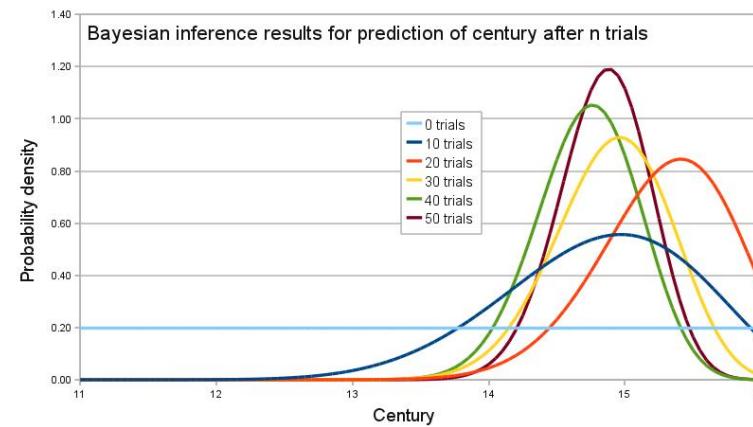


Introduction

- Why Probabilistic Models:
 - Overconfident point-estimate models
 - Interpretability and explainability
- Probability and Bayesian Inference:
 - Probability distributions and sampling
 - Start with a prior belief
 - Provide/sample new data
 - Update (variational) posterior belief



Source: [Reddit](#)

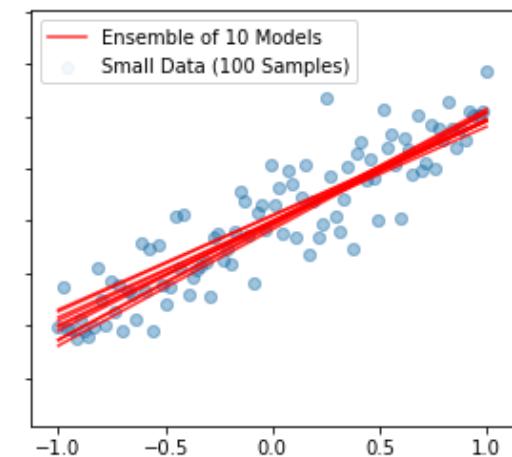
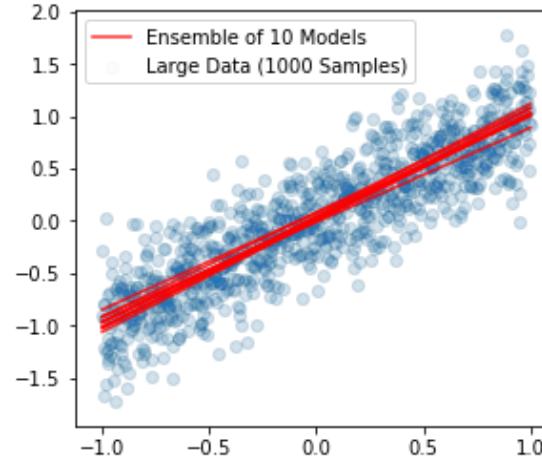
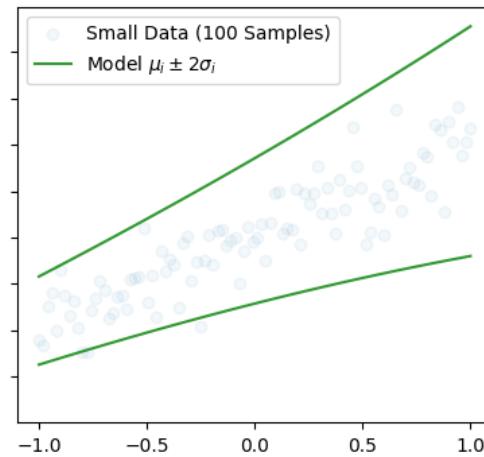
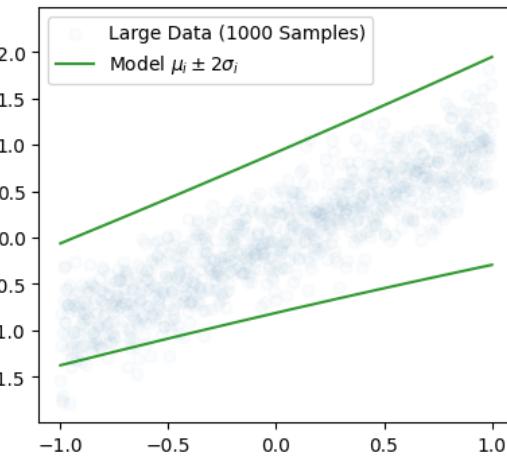


Source: [Wikipedia](#)



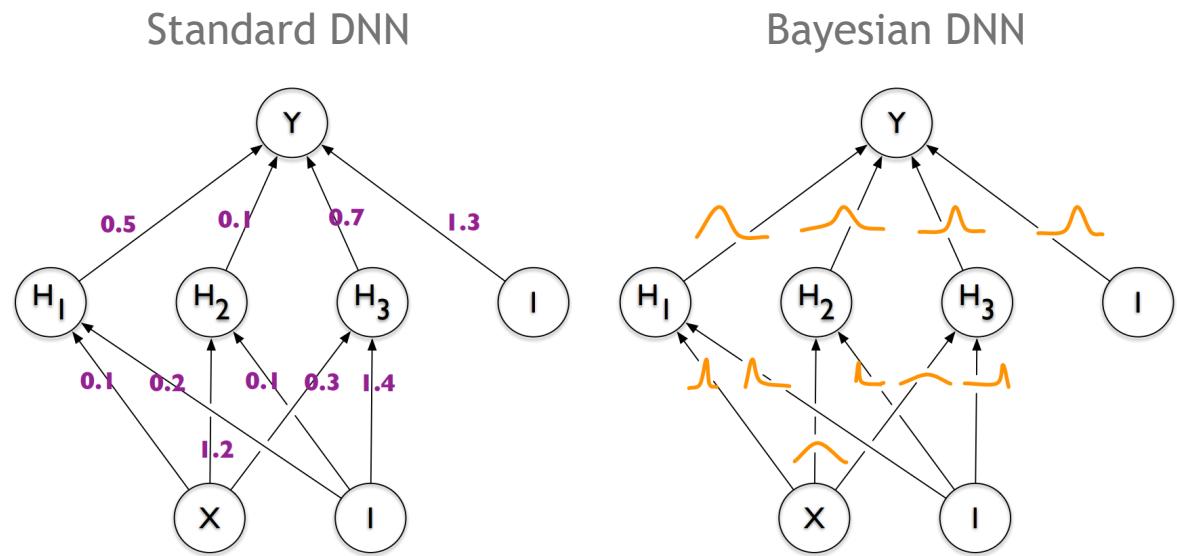
Uncertainty Quantification

- Aleatoric Uncertainty
 - Homoscedastic
 - Heteroscedastic
- Epistemic Uncertainty



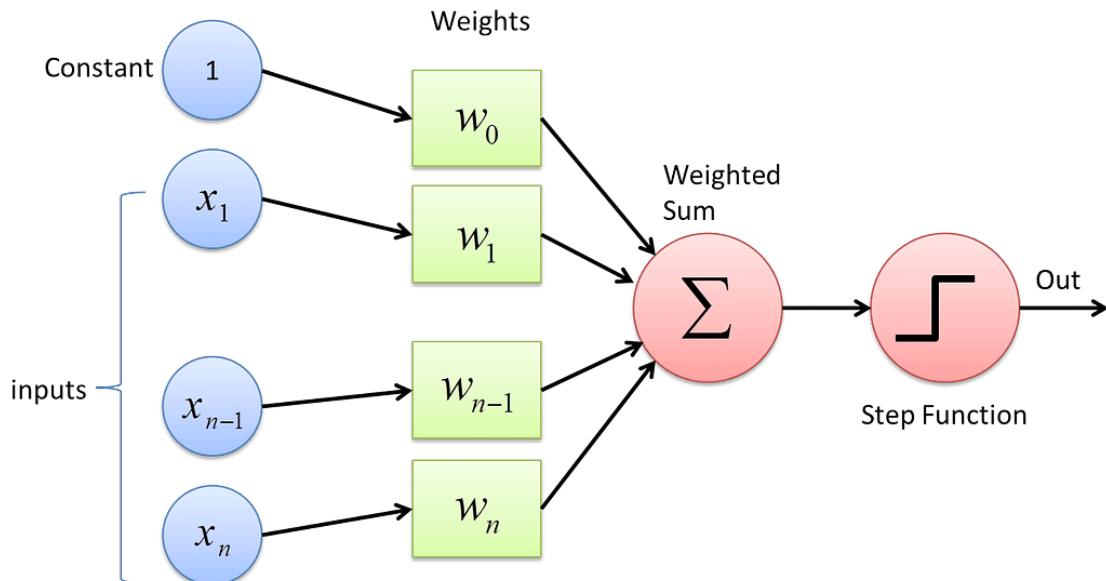
What is Probabilistic Deep Learning?

- Standard (Point-Estimate) Neural Networks:
 - Given input, predict output
 - Learn θ , where $w, b \in \theta$
- Probabilistic Neural Networks:
 - Given input, predict likelihood of output
 - Learn PDF $\sim \theta$, where $w, b \in \theta$
- Bayesian Neural Networks:
 - Given input, predict output with certainty
 - Learn μ, Σ where $w_i, b_i \sim \mathcal{N}(\mu_i, \sigma_i)$
- Probabilistic Bayesian Neural Networks:
 - PNN + BNN



Source: [Charles, et al.](#)

Mechanics of Bayesian Deep Learning



Source: [Sharma, TDS](#)

- Notation:

Input: x	Output: \hat{y}	Weights: w	Bias: b
Data: $D = (x_i, y_i)_{i=1}^n$		DNN Model: $F_\theta()$	
		Activation Function: $g()$	Loss Function: $\mathcal{L}()$

- Forward-Propagation:

$$\begin{aligned}\hat{y} &= g(w^T x + b) \\ &= g(w_1 x_1 + w_2 x_2 + w_3 x_3 + b)\end{aligned}$$

Neural Network

$$w_i, b_i \in \mathbb{R}$$

To learn $\{w_i, b_i\}$

Bayesian Neural Network

$$w_i, b_i \sim \mathcal{N}(\mu_i, \sigma_i)$$

To learn $\{\mu_i, \sigma_i\}$

Mechanics of Bayesian Deep Learning

Neural Network

- Training:

$$\theta^* = \underset{\theta}{\operatorname{argmax}} \sum_{(x_i, y_i) \in D} \log[p(y_i | x_i, \theta)]$$

$$\theta^* = \underset{\theta}{\operatorname{argmin}} \sum_{(x_i, y_i) \in D} \mathcal{L}(F_\theta(x_i), y_i)$$

- Prediction:

$$p(\hat{y} | \hat{x}, \theta^*)$$

$$\hat{y} = F_{\theta^*}(\hat{x})$$

Bayesian Neural Network

$$\mu^*, \Sigma^* = \underset{\mu, \Sigma}{\operatorname{argmax}} \sum_{(x_i, y_i) \in D} \log[p(y_i | x_i, \theta)] - KL[p(\theta), p(\theta_0)]$$

$$\theta \sim \mathcal{N}(\mu, \Sigma) \quad \theta_0 \sim \mathcal{N}(0, I)$$

$$\mu^*, \Sigma^* = \underset{\mu, \Sigma}{\operatorname{argmin}} \sum_{(x_i, y_i) \in D} \mathcal{L}(F_\theta(x_i), y_i) + KL[p(\theta), p(\theta_0)]$$

$$p(\hat{y} | \hat{x}, D) = \int p(\hat{y} | \hat{x}, \theta^*) p(\theta^* | D) d\theta \quad \theta^* \sim \mathcal{N}(\mu^*, \Sigma^*)$$

$$\hat{y} = \frac{1}{K} \sum_{k=1}^K F_{\theta_k^*}(\hat{x}) \quad \theta_k^* \sim \mathcal{N}(\mu^*, \Sigma^*)$$



Mechanics of Bayesian Deep Learning

- Back-Propagation:

- Intractability: must calculate derivatives of the parameters being learnt, i.e., derivatives of distributions

$$p(\hat{y} | \hat{x}, D) = \int p(\hat{y} | \hat{x}, \theta^*) p(\theta^* | D) d\theta$$

- Solution: local re-parameterization trick -- “moves” the parameters to be learnt (μ, σ in case of a Gaussian distribution), out of the distribution function for any weight w .

w.k.t. $\theta = (\mu, \sigma)$

let $\epsilon \sim \mathcal{N}(0, 1)$

$f(\epsilon) = w = \mu + \sigma \cdot \epsilon$

compute derivatives

- Parameter Updates:

$$\Delta\mu = \frac{\partial f}{\partial w} + \frac{\partial f}{\partial \mu} \quad \Delta\sigma = \frac{\partial f}{\partial w} \frac{\epsilon}{\sigma} + \frac{\partial f}{\partial \sigma}$$

$$\mu := \mu - \alpha \Delta\mu \quad \sigma := \sigma - \alpha \Delta\sigma$$

$$\theta^* = (\mu^*, \sigma^*)$$



Pros and Cons of Bayesian Deep Learning

- Pros

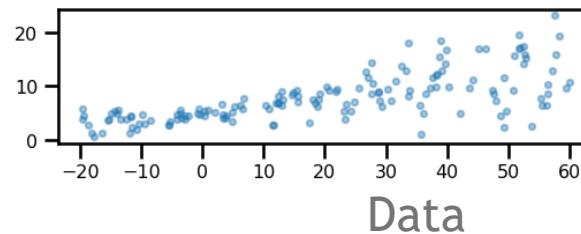
- They are more robust and generalizable
- They can quantify the uncertainty in their predictive output

- Cons

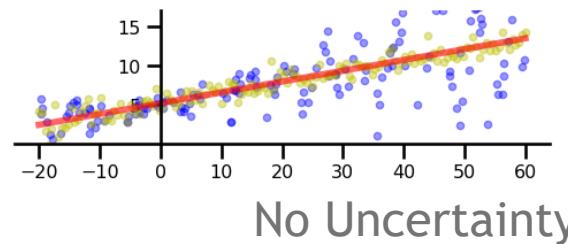
- They can be more complicated to train (may require knowledge of probability and statistics)
- They can be slower to converge and require more data (since the weights of the network are distributions instead of single values)



Preliminary Linear Regression Experiments

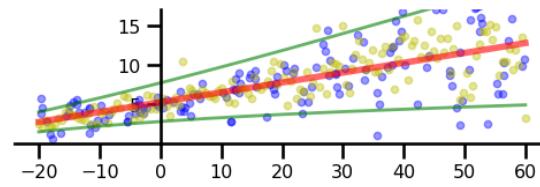


Training Data



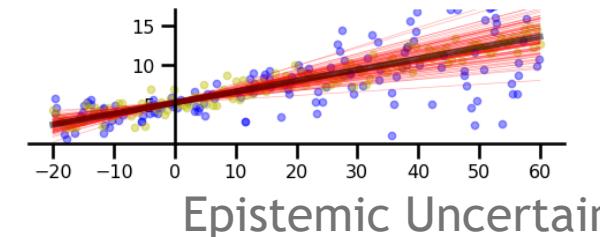
No Uncertainty

Train Data Samples
Test Data Predictions
Prediction μ



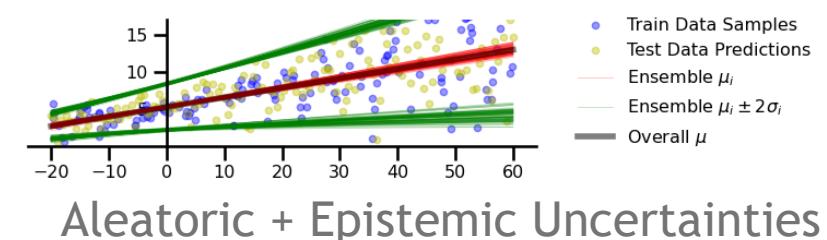
Aleatoric Uncertainty

Train Data Samples
Test Data Predictions
Prediction μ
Prediction $\mu \pm 2\sigma$



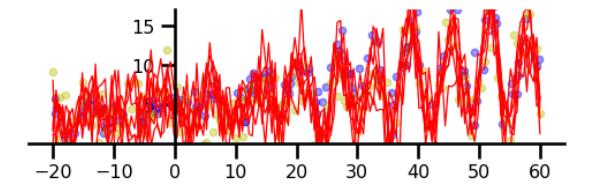
Epistemic Uncertainty

Train Data Samples
Test Data Predictions
Ensemble μ_i
Overall μ



Aleatoric + Epistemic Uncertainties

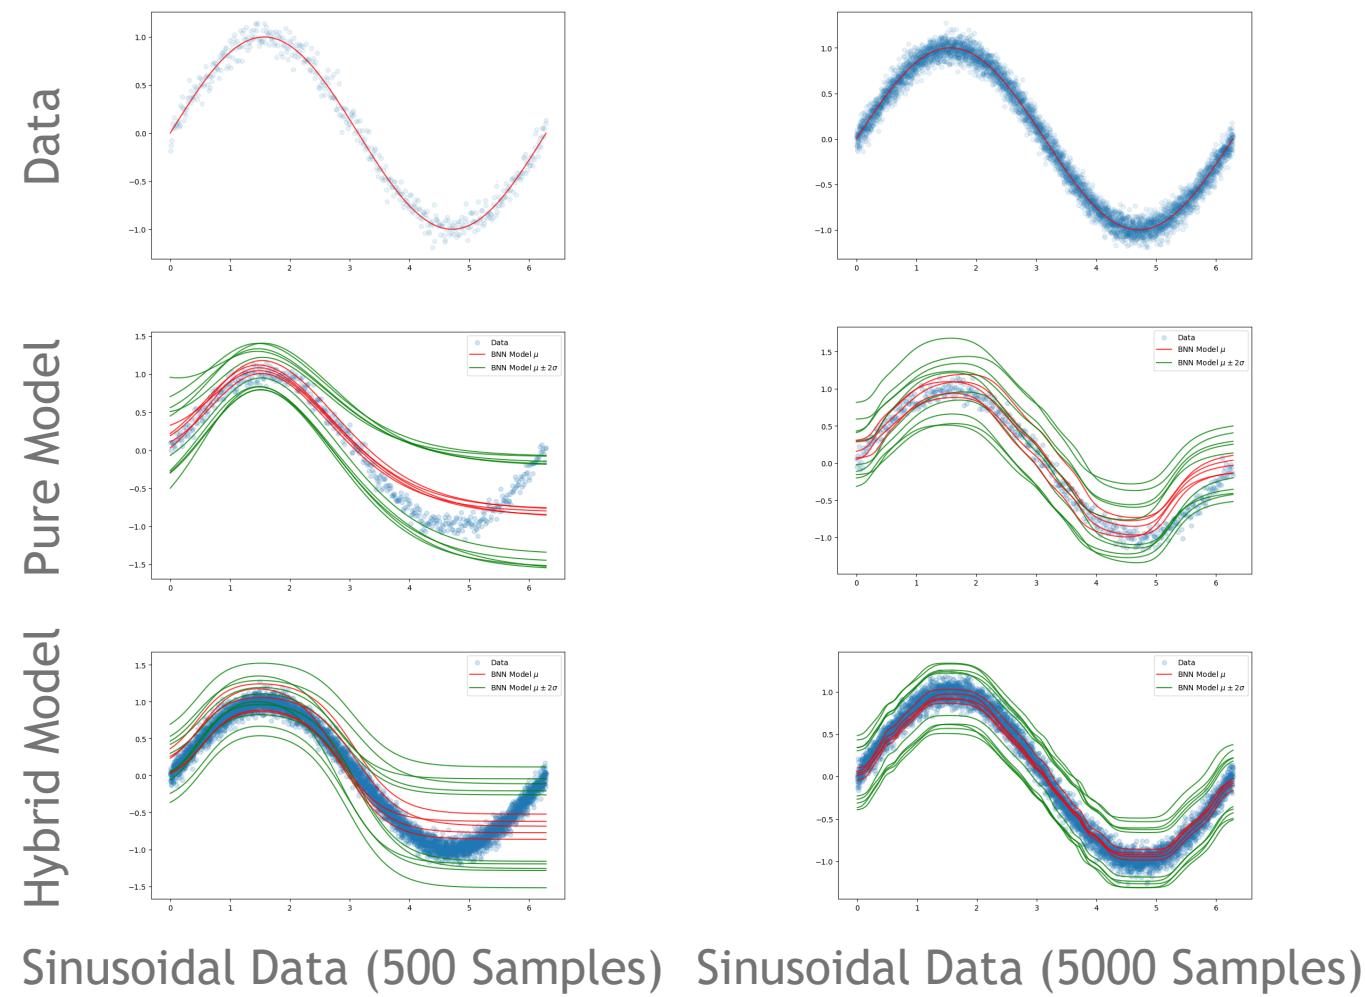
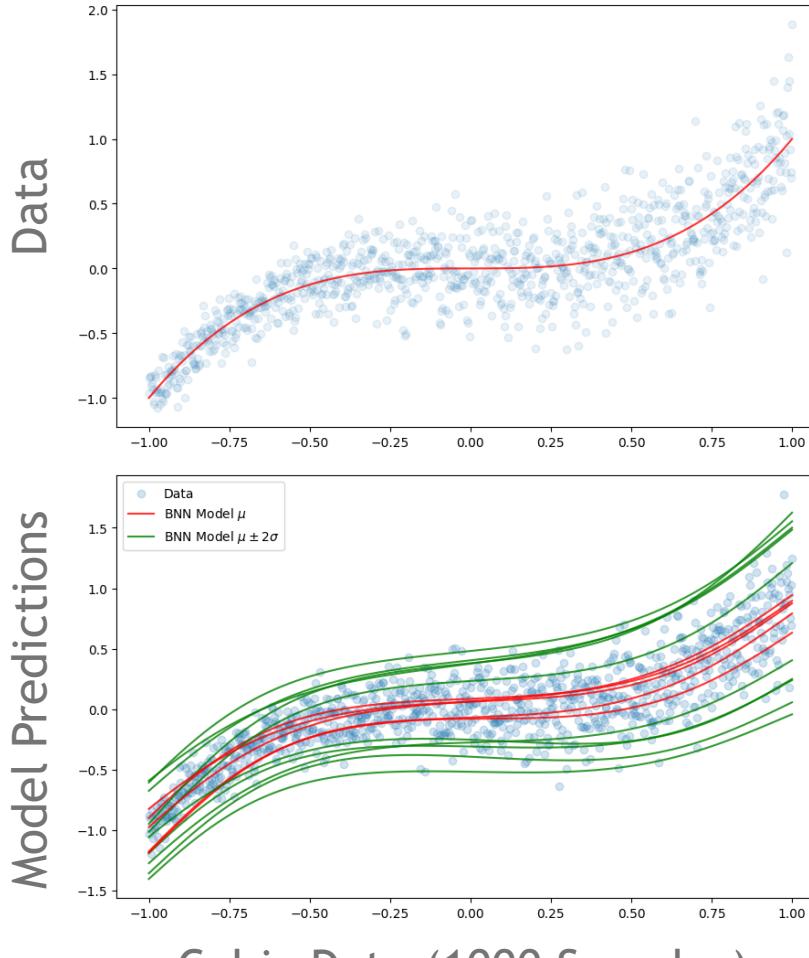
Train Data Samples
Test Data Predictions
Ensemble μ_i
Ensemble $\mu_i \pm 2\sigma_i$
Overall μ



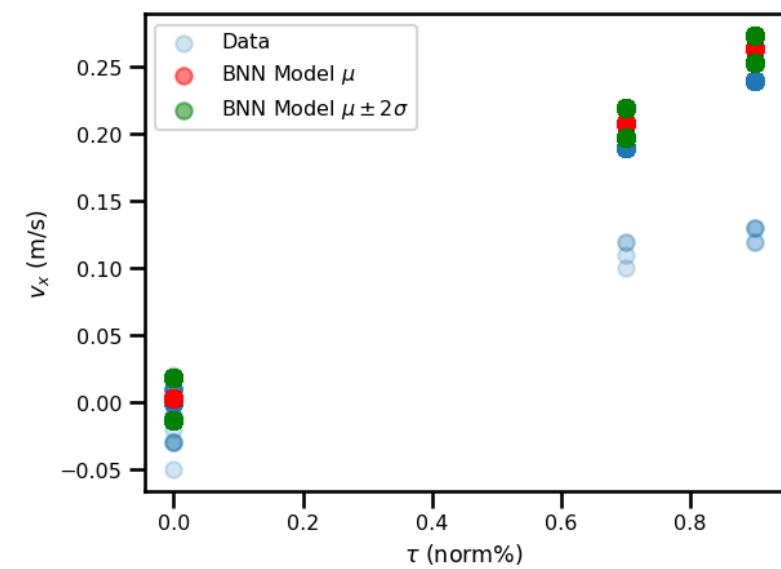
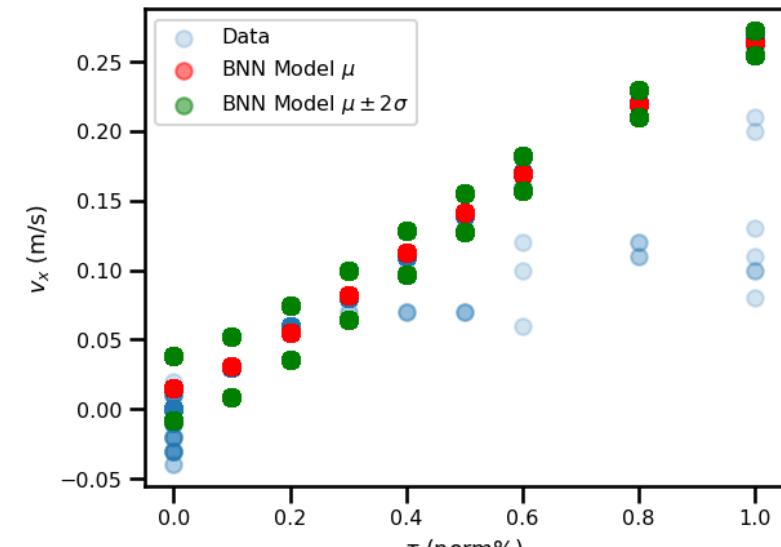
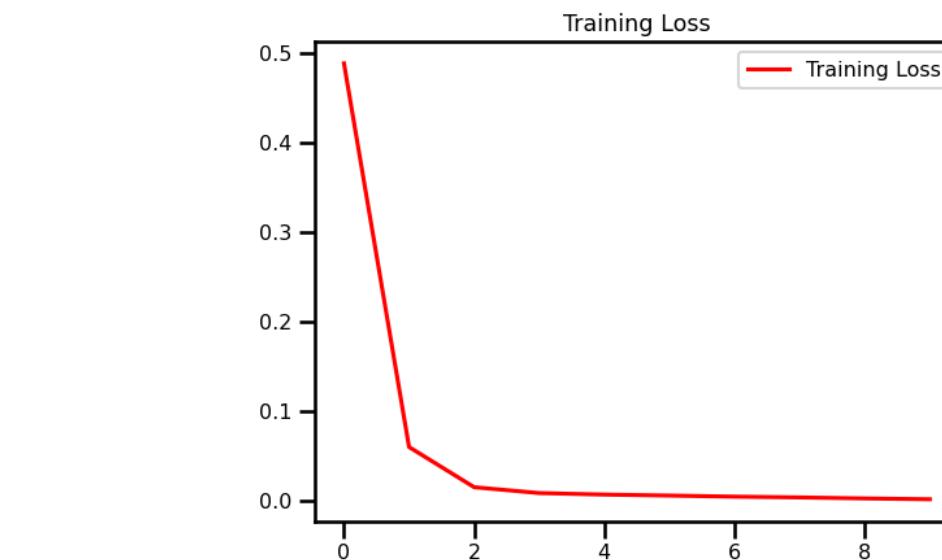
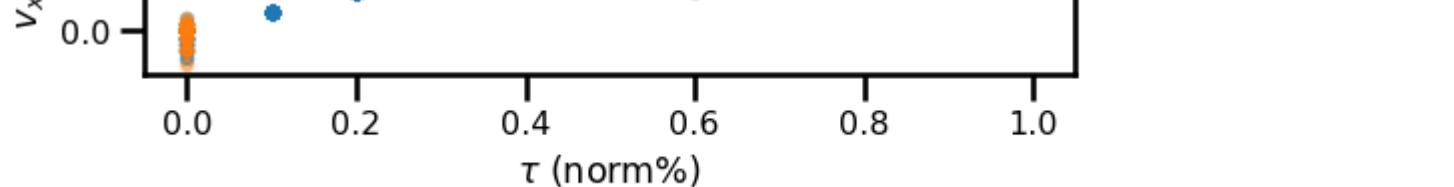
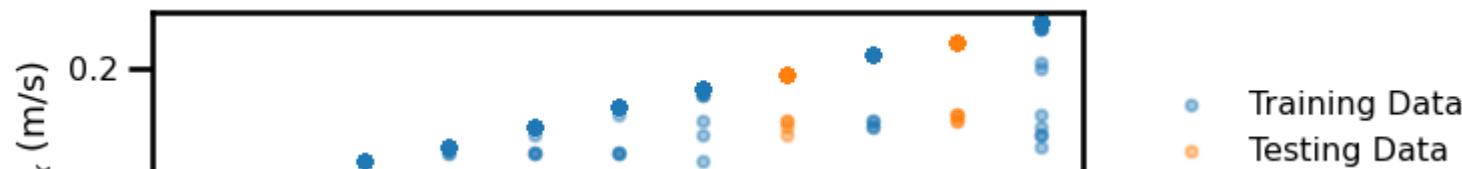
Functional Uncertainty

Train Data Samples
Test Data Predictions
Ensemble μ_i

Preliminary Nonlinear Regression Experiments

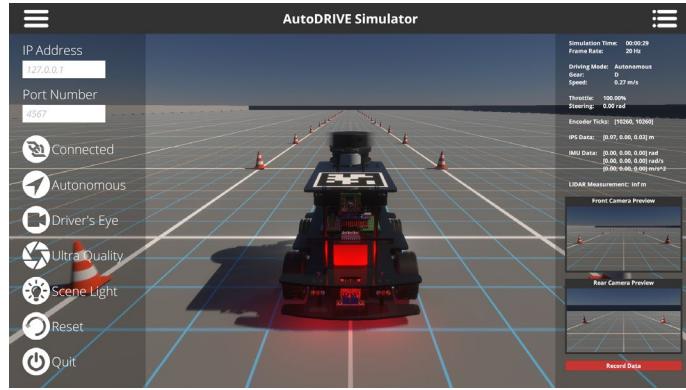


Preliminary 1D VD Experiments



Data Collection

- AutoDRIVE Simulator
- Nigel (1:14) Vehicle
- Dynamics + perception
- Automated script



DATA	timestamp	throttle	steering	leftTicks	rightTicks	posX	posY	posZ	roll	pitch	yaw	speed	angX	angY	angZ	accX	accY	accZ	cam0	cam1
UNIT	yyyy_MM_dd_HH_mm_ss_fff	norm%	rad	count	count	m	m	m	rad	rad	rad	m/s	rad/s	rad/s	rad/s	m/s^2	m/s^2	m/s^2	img_path	img_path

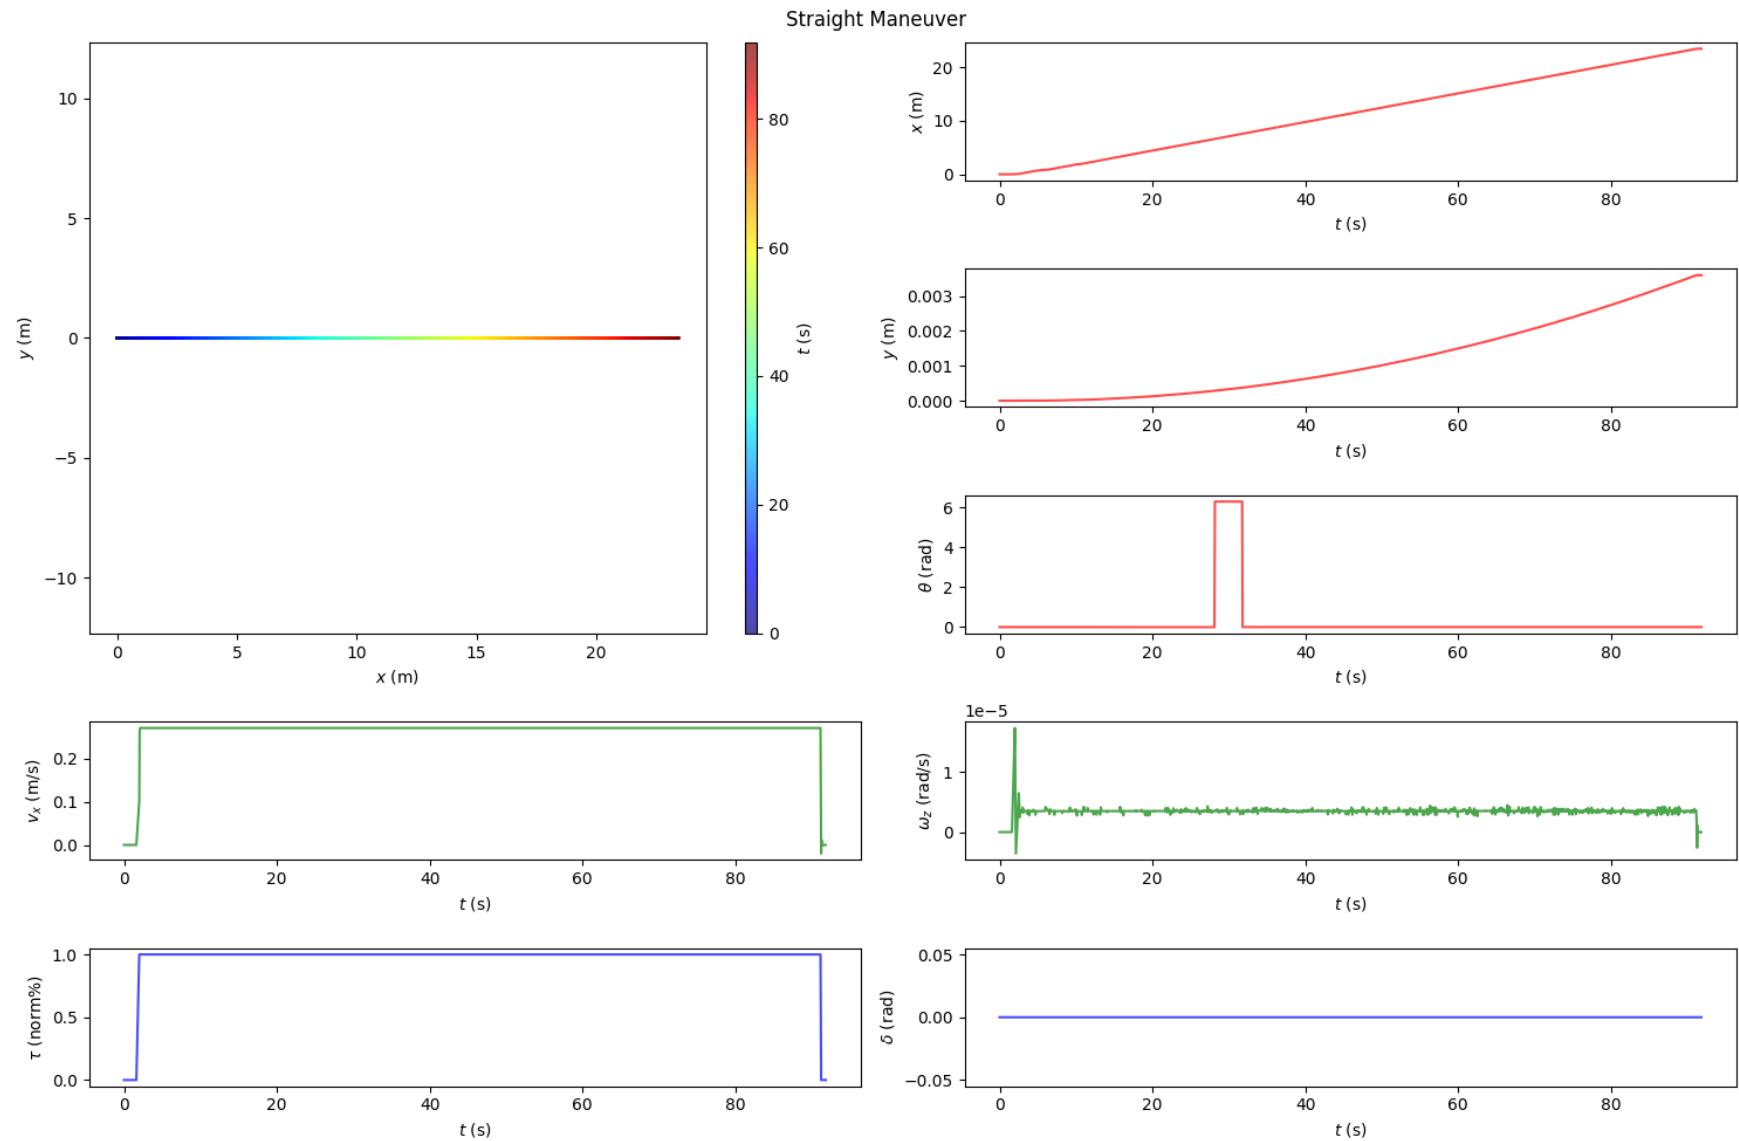
DATA	timestamp	state
UNIT	yyyy_MM_dd_HH_mm_ss_fff	Int{0,1,2,3}

Dataset: <https://github.com/Tinker-Twins/AutoDRIVE-Nigel-Dataset>



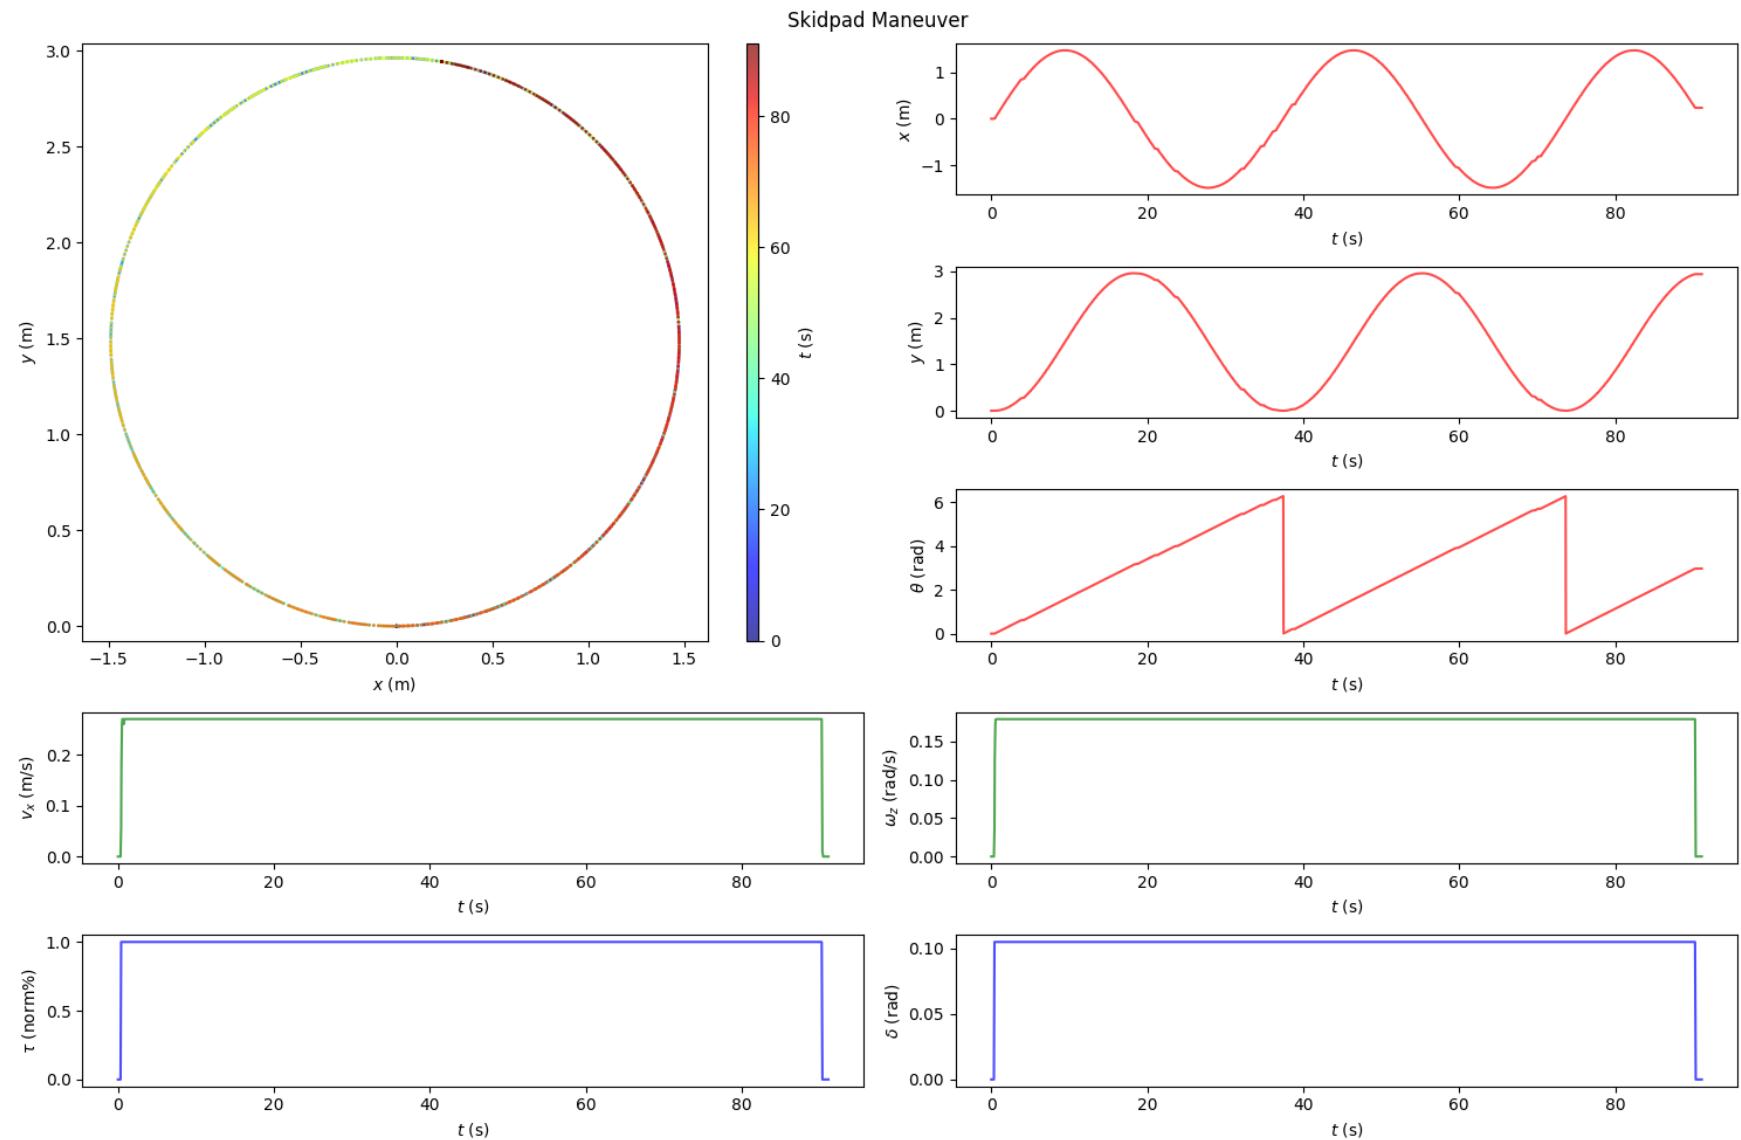
Data Collection

- Straight
- Skidpad
- Fishhook
- Slalom
- Eight



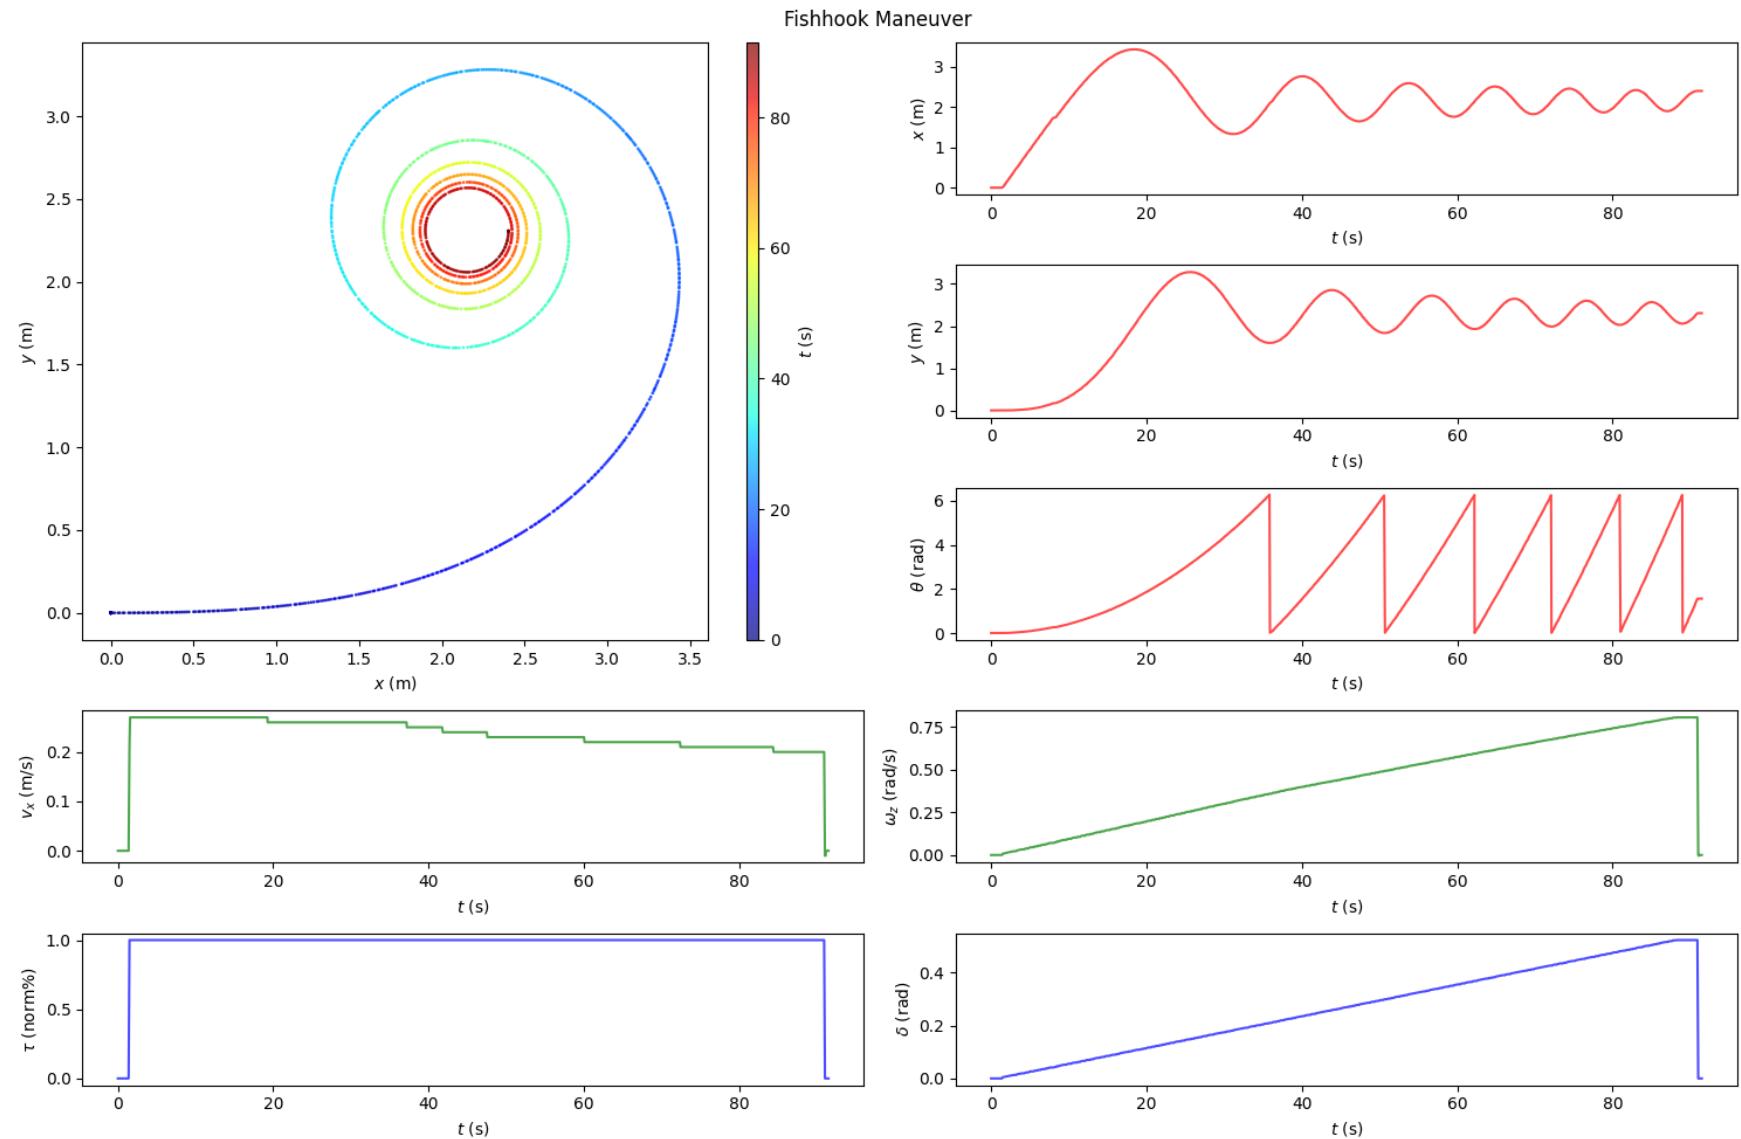
Data Collection

- Straight
- Skidpad
- Fishhook
- Slalom
- Eight



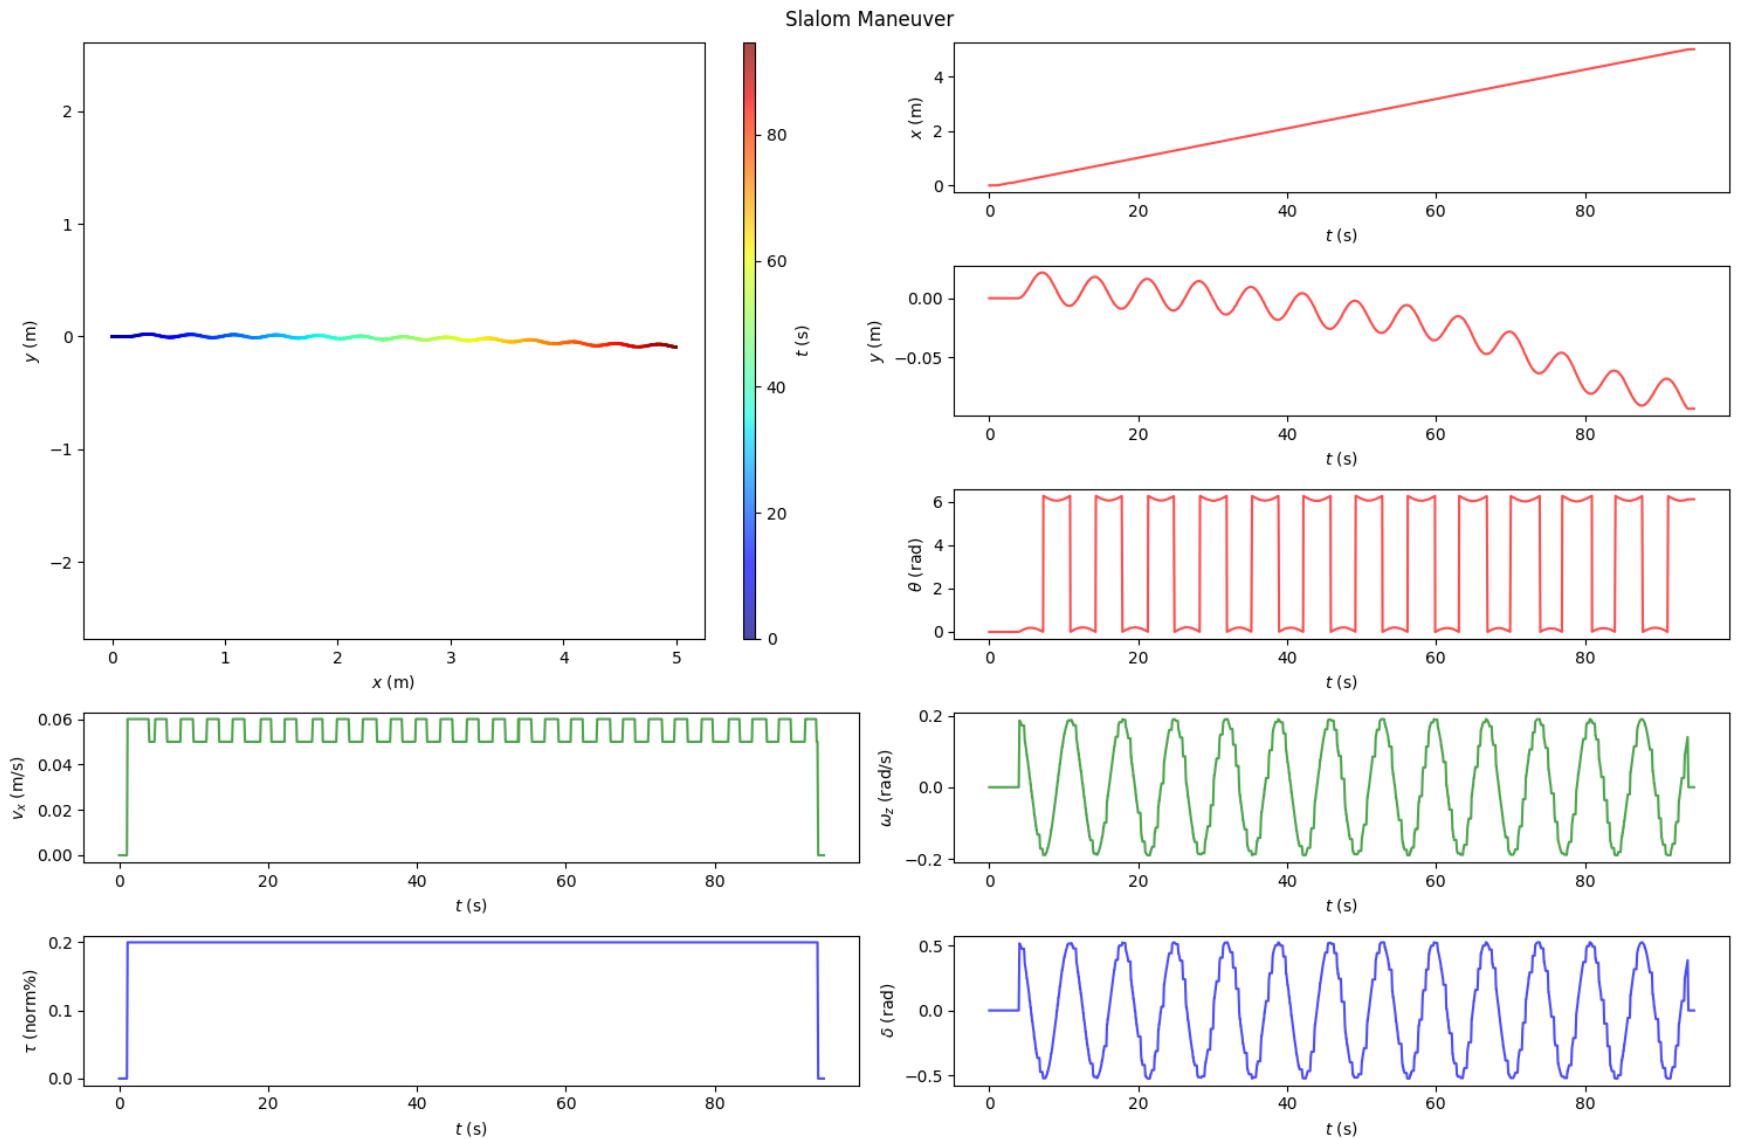
Data Collection

- Straight
- Skidpad
- **Fishhook**
- Slalom
- Eight



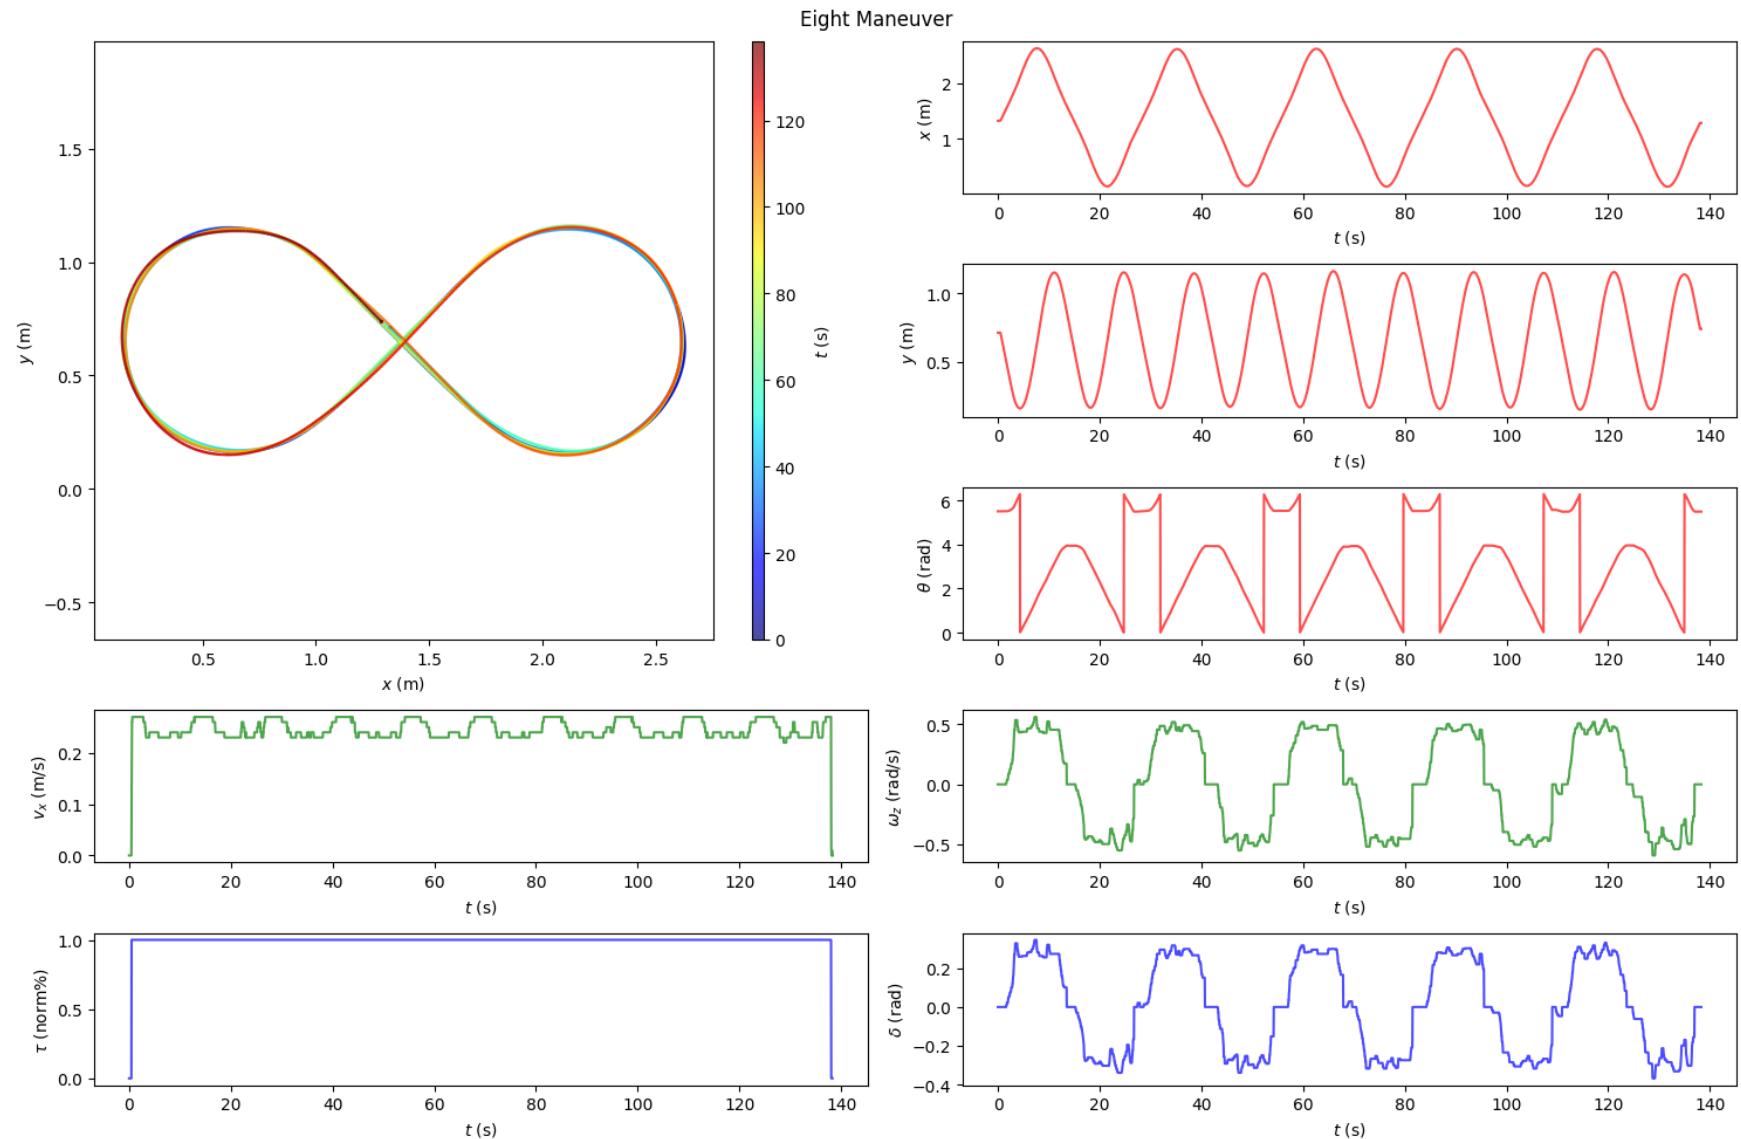
Data Collection

- Straight
- Skidpad
- Fishhook
- **Slalom**
- Eight



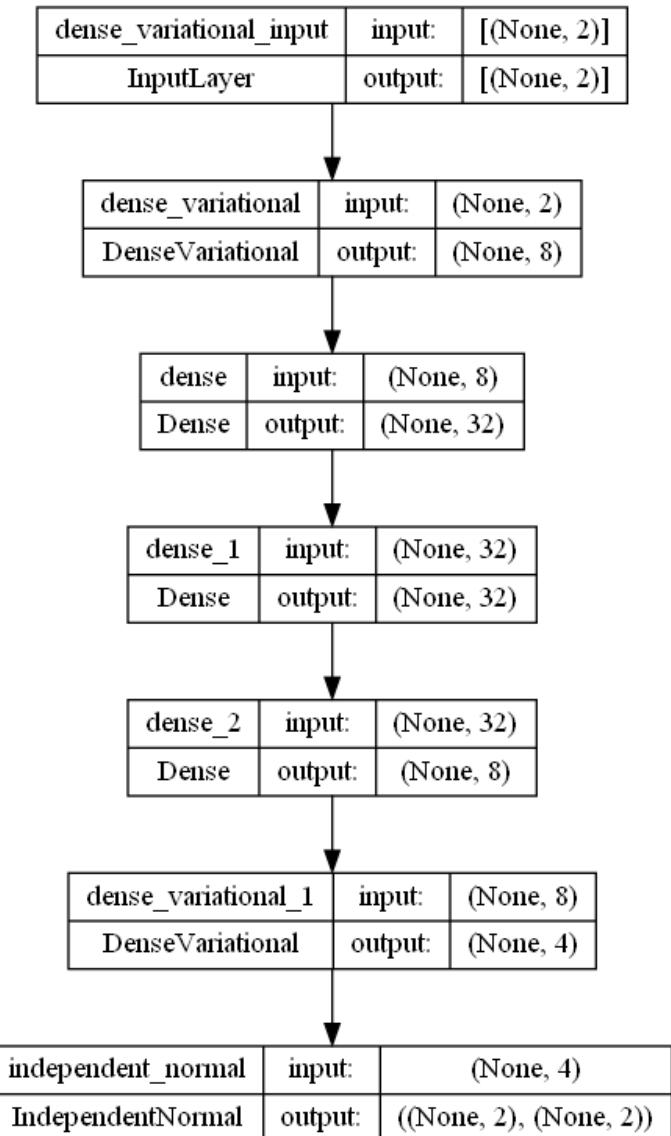
Data Collection

- Straight
- Skidpad
- Fishhook
- Slalom
- Eight



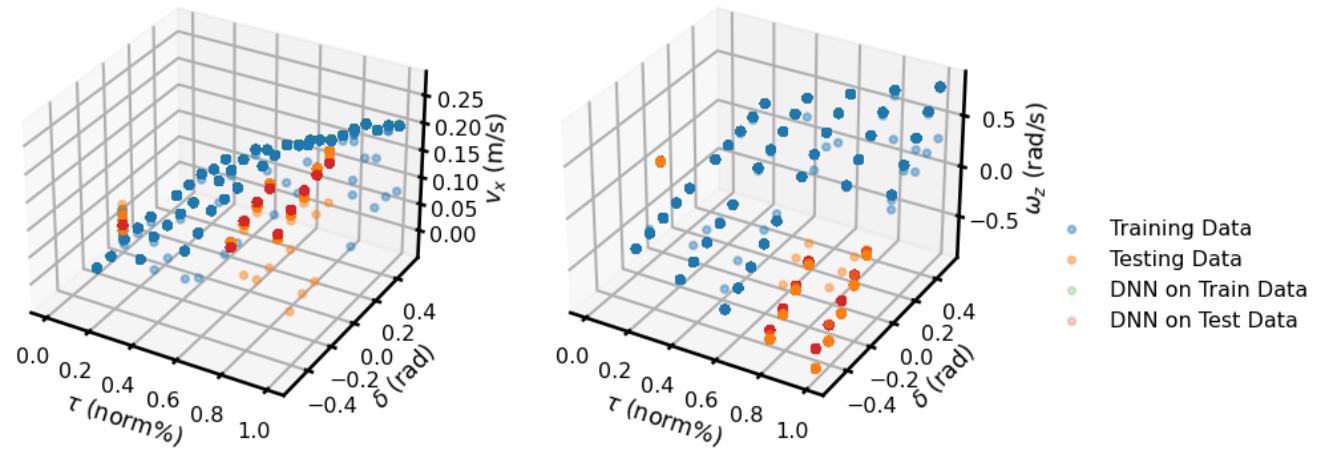
Hyperparameter Tuning

- # epochs: {5, 10, 15, 20, **25**, 30}
- # hidden layers: {3d, 2v, 3v, 5v, **3d2v**}
- Learning rate: {5e-3, **1e-3**, 1e-2, 1e-1}
- Activation functions:
 - Bounded: {tanh, **sigmoid**}
 - Unbounded: {elu, **relu**, selu, softplus}
- Optimizer: {SGD, **RMSprop**, Adam}
- Loss function: {NLL, **MSE**, MAE}
- Prior distribution: {univariate Gaussian, **multivariate Gaussian**}
- Posterior distribution: {univariate Gaussian, **multivariate Gaussian**}
- KL-divergence: {approx, **exact**}



Hyperparameter Tuning

- # epochs: {5, 10, 15, 20, **25**, 30}
- # hidden layers: {3d, 2v, 3v, 5v, **3d2v**}
- Learning rate: {5e-3, **1e-3**, 1e-2, 1e-1}
- Activation functions:
 - Bounded: {tanh, **sigmoid**}
 - Unbounded: {elu, **relu**, selu, softplus}
- Optimizer: {SGD, **RMSprop**, Adam}
- Loss function: {NLL, **MSE**, MAE}
- Prior distribution: {univariate Gaussian, **multivariate Gaussian**}
- Posterior distribution: {univariate Gaussian, **multivariate Gaussian**}
- KL-divergence: {approx, **exact**}

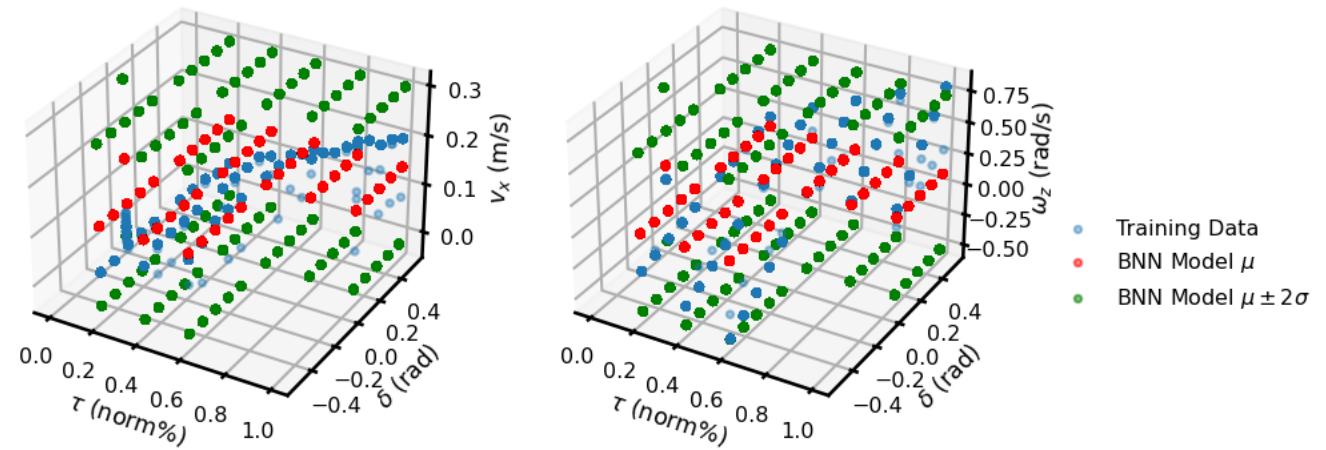


Skidpad {5, 1e-3, 3d, relu, Adam, MSE}



Hyperparameter Tuning

- # epochs: {5, 10, 15, 20, **25**, 30}
- # hidden layers: {3d, 2v, 3v, 5v, **3d2v**}
- Learning rate: {5e-3, **1e-3**, 1e-2, 1e-1}
- Activation functions:
 - Bounded: {tanh, **sigmoid**}
 - Unbounded: {elu, **relu**, selu, softplus}
- Optimizer: {SGD, **RMSprop**, Adam}
- Loss function: {NLL, **MSE**, MAE}
- Prior distribution: {univariate Gaussian, **multivariate Gaussian**}
- Posterior distribution: {univariate Gaussian, **multivariate Gaussian**}
- KL-divergence: {approx, **exact**}

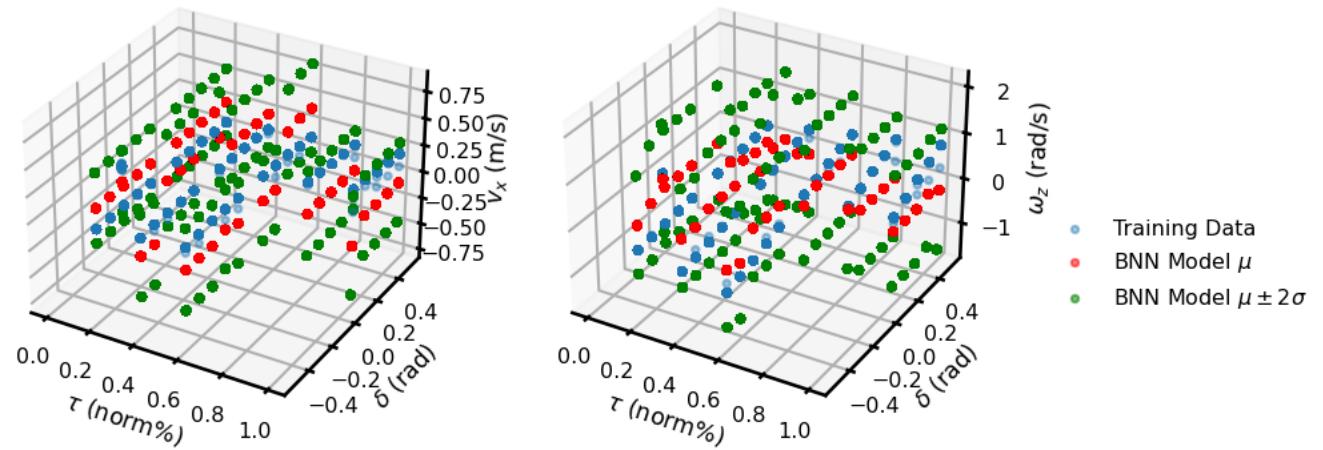


Skidpad {10, 1e-3, 3v, sigmoid, Adam, NLL, MVG, approx KLD}



Hyperparameter Tuning

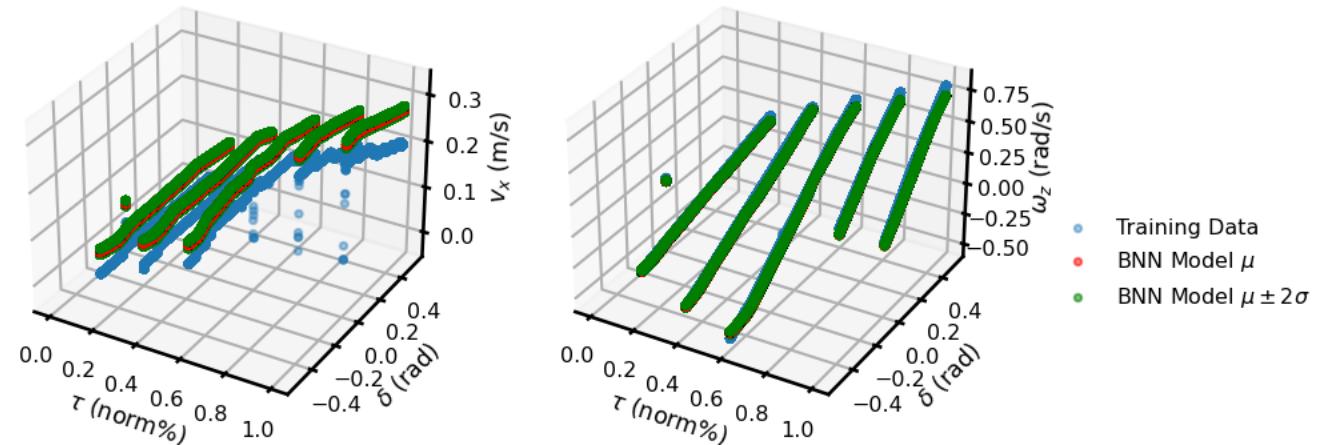
- # epochs: {5, 10, 15, 20, **25**, 30}
- # hidden layers: {3d, 2v, 3v, 5v, **3d2v**}
- Learning rate: {5e-3, **1e-3**, 1e-2, 1e-1}
- Activation functions:
 - Bounded: {tanh, **sigmoid**}
 - Unbounded: {elu, **relu**, selu, softplus}
- Optimizer: {SGD, **RMSprop**, Adam}
- Loss function: {NLL, **MSE**, MAE}
- Prior distribution: {univariate Gaussian, **multivariate Gaussian**}
- Posterior distribution: {univariate Gaussian, **multivariate Gaussian**}
- KL-divergence: {approx, **exact**}



Skidpad {25, 5e-3, 2v, tanh, RMSProp, NLL, MVG, approx KLD}

Hyperparameter Tuning

- # epochs: {5, 10, 15, 20, **25**, 30}
- # hidden layers: {3d, 2v, 3v, 5v, **3d2v**}
- Learning rate: {5e-3, **1e-3**, 1e-2, 1e-1}
- Activation functions:
 - Bounded: {tanh, **sigmoid**}
 - Unbounded: {elu, **relu**, selu, softplus}
- Optimizer: {SGD, **RMSprop**, Adam}
- Loss function: {NLL, **MSE**, MAE}
- Prior distribution: {univariate Gaussian, **multivariate Gaussian**}
- Posterior distribution: {univariate Gaussian, **multivariate Gaussian**}
- KL-divergence: {approx, **exact**}

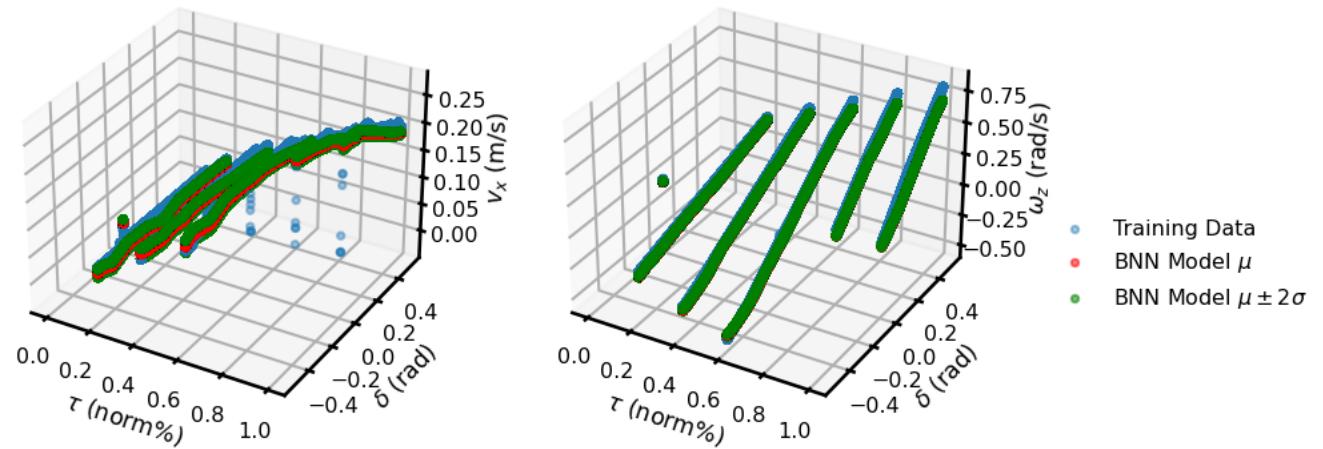


Fishhook {25, 1e-3, 3d2v, sigmoid, RMSProp, MAE, MVG, approx KLD}



Hyperparameter Tuning

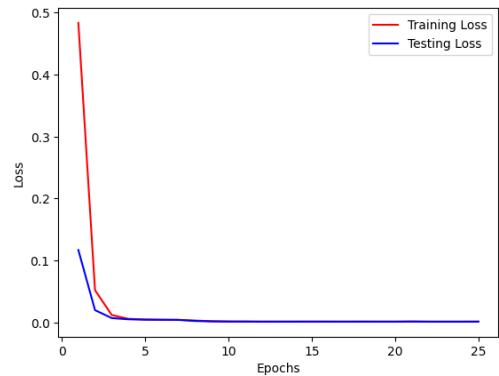
- # epochs: {5, 10, 15, 20, **25**, 30}
- # hidden layers: {3d, 2v, 3v, 5v, **3d2v**}
- Learning rate: {5e-3, **1e-3**, 1e-2, 1e-1}
- Activation functions:
 - Bounded: {tanh, **sigmoid**}
 - Unbounded: {elu, **relu**, selu, softplus}
- Optimizer: {SGD, **RMSprop**, Adam}
- Loss function: {NLL, **MSE**, MAE}
- Prior distribution: {univariate Gaussian, **multivariate Gaussian**}
- Posterior distribution: {univariate Gaussian, **multivariate Gaussian**}
- KL-divergence: {approx, **exact**}



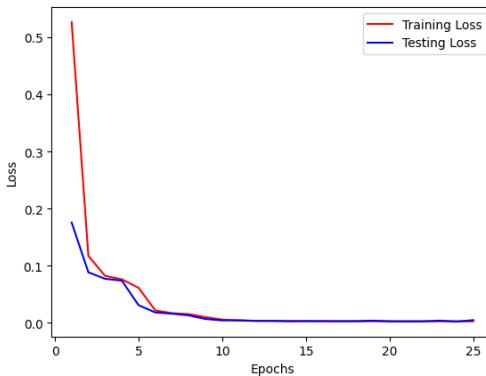
Fishhook {25, 1e-3, 3d2v, sigmoid, Adam, MAE, MVG, exact KLD}



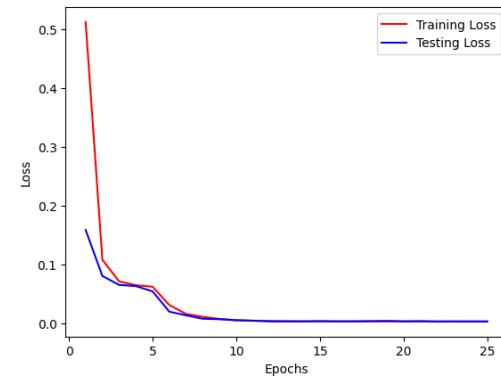
Model Training



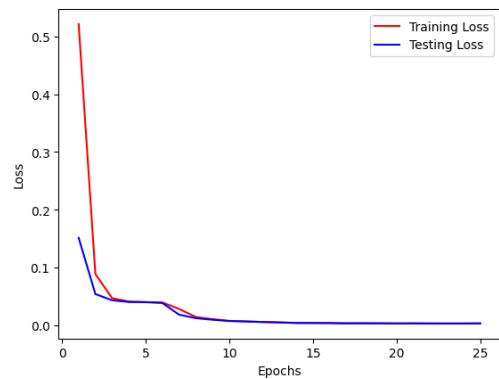
Straight Maneuver



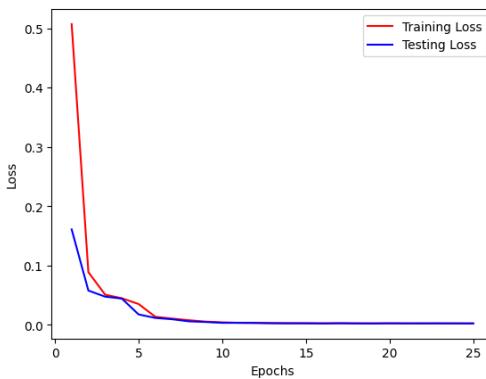
Skidpad Maneuver



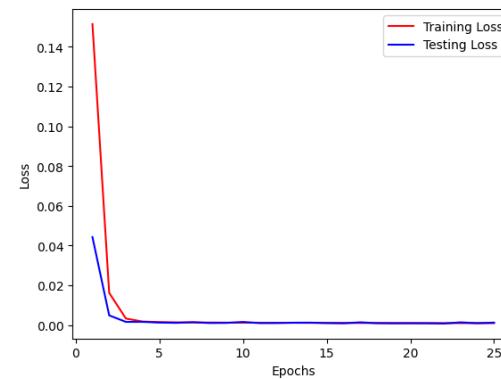
Fishhook Maneuver



Slalom Maneuver



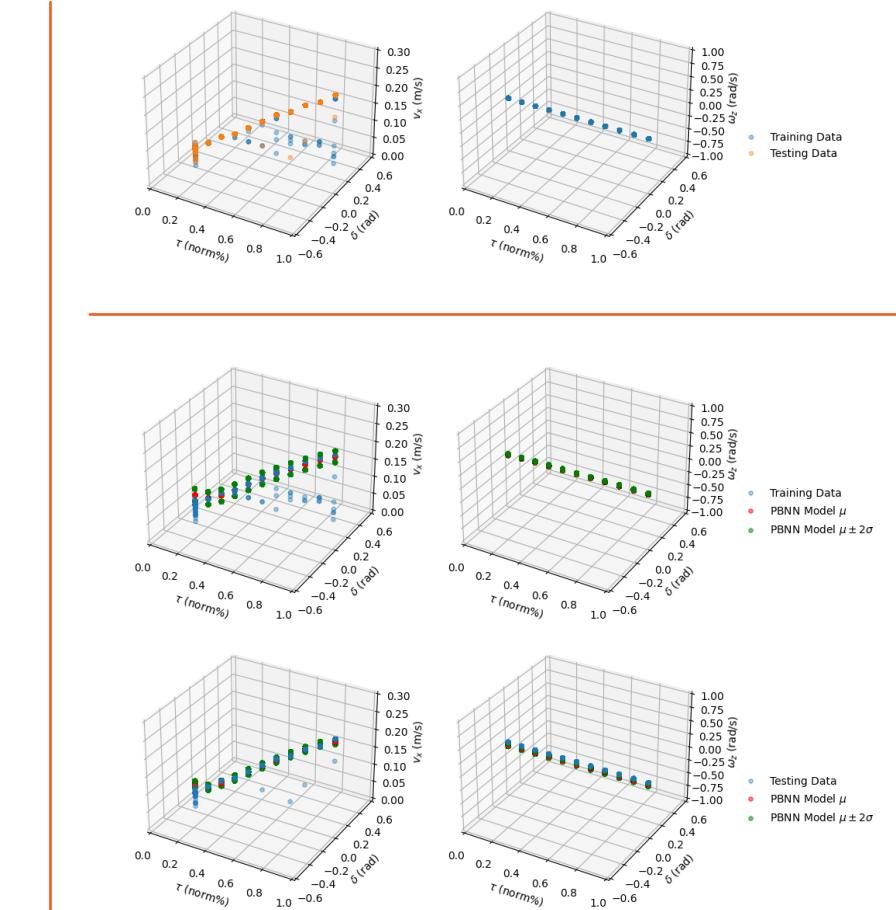
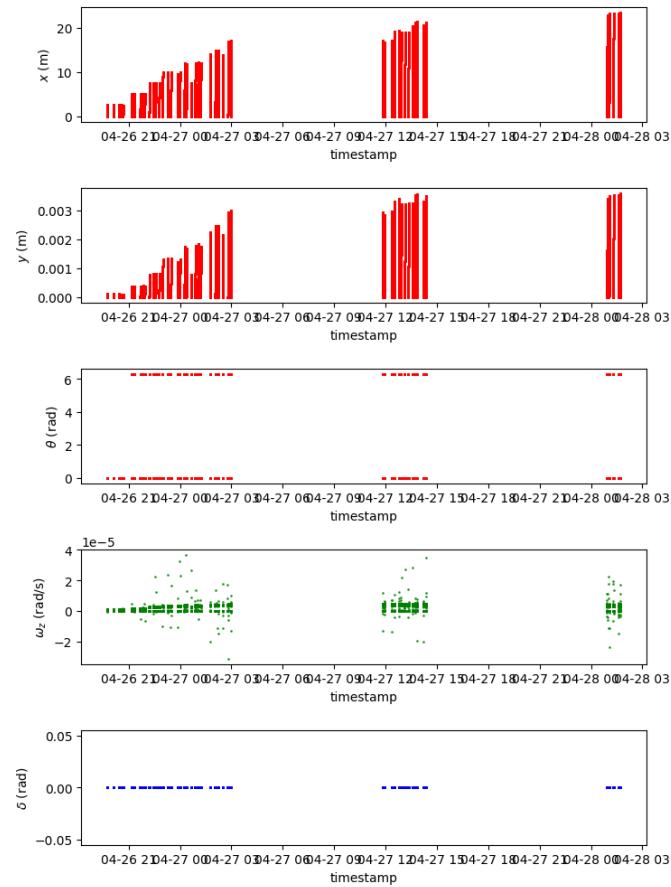
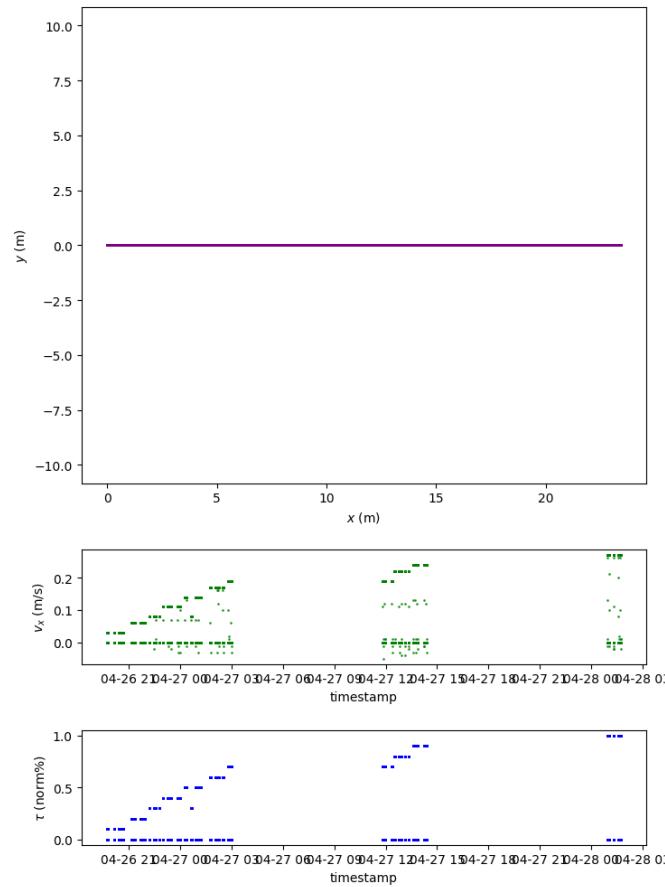
Eight Maneuver



All Maneuvers



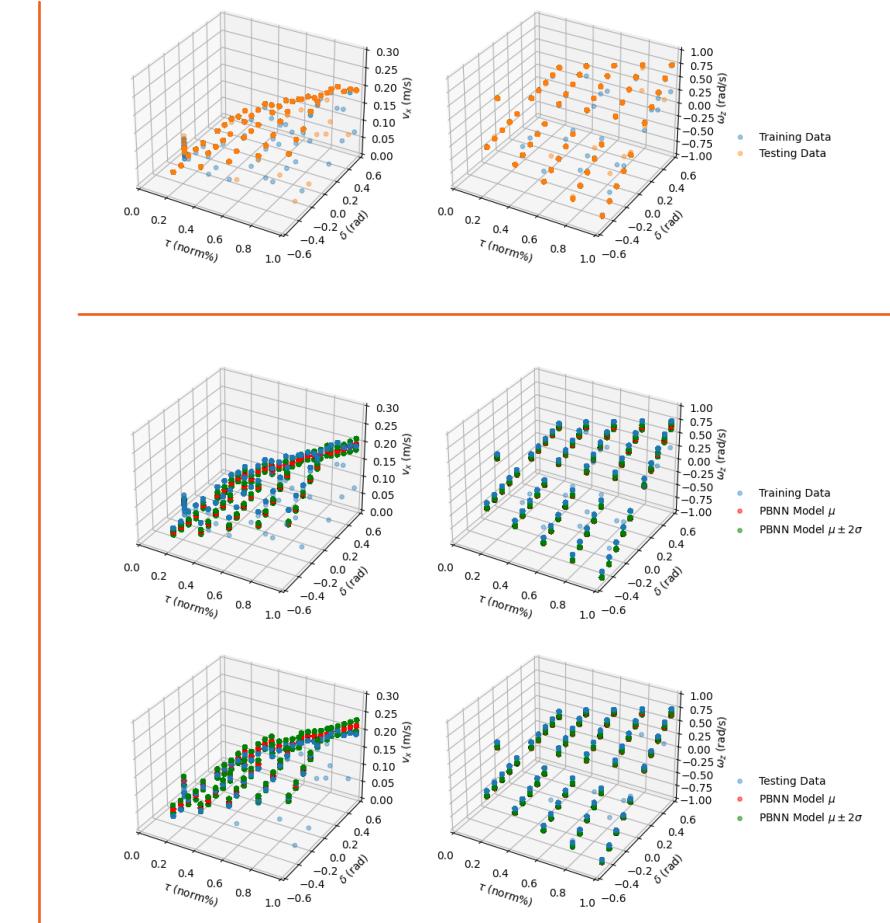
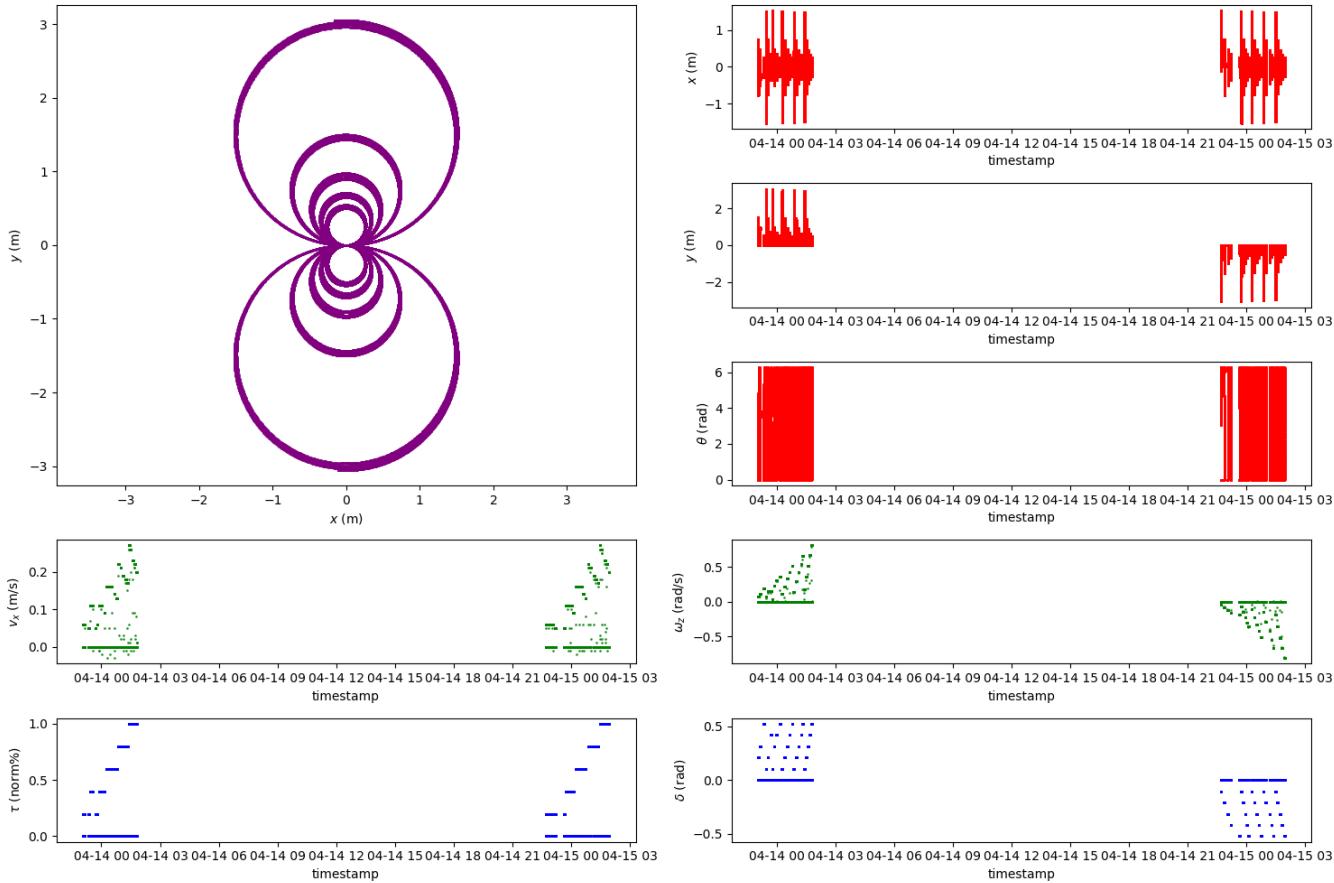
Model Inference - Straight Maneuver



Note: This work adheres to Markov assumption.



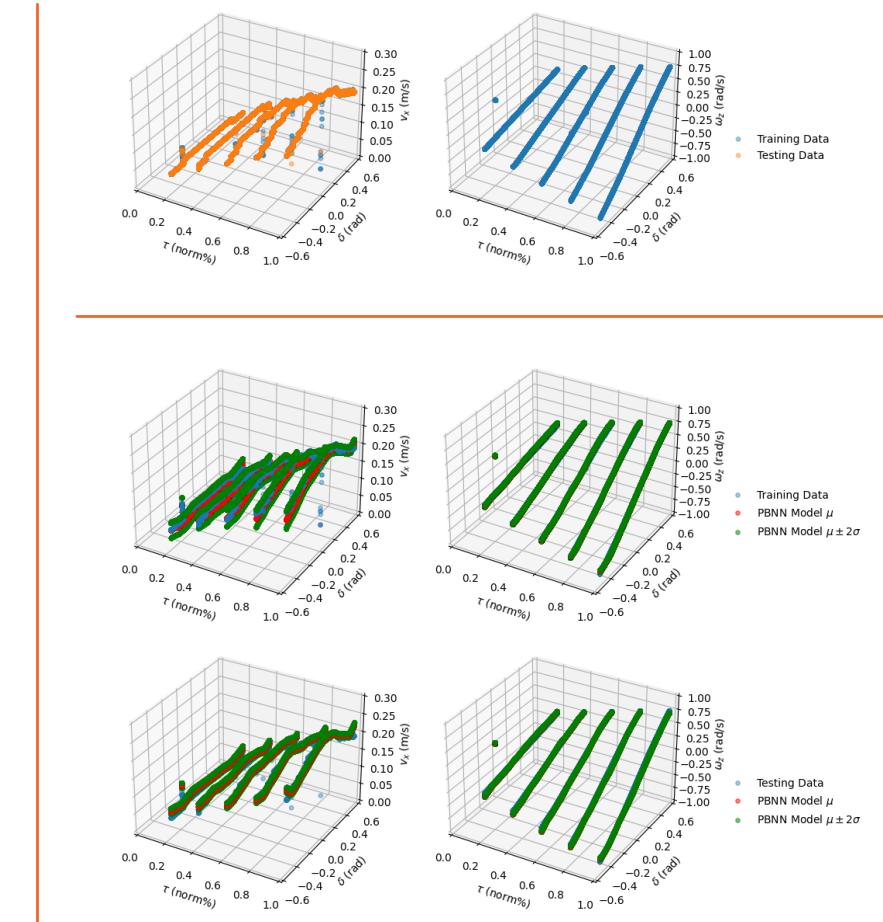
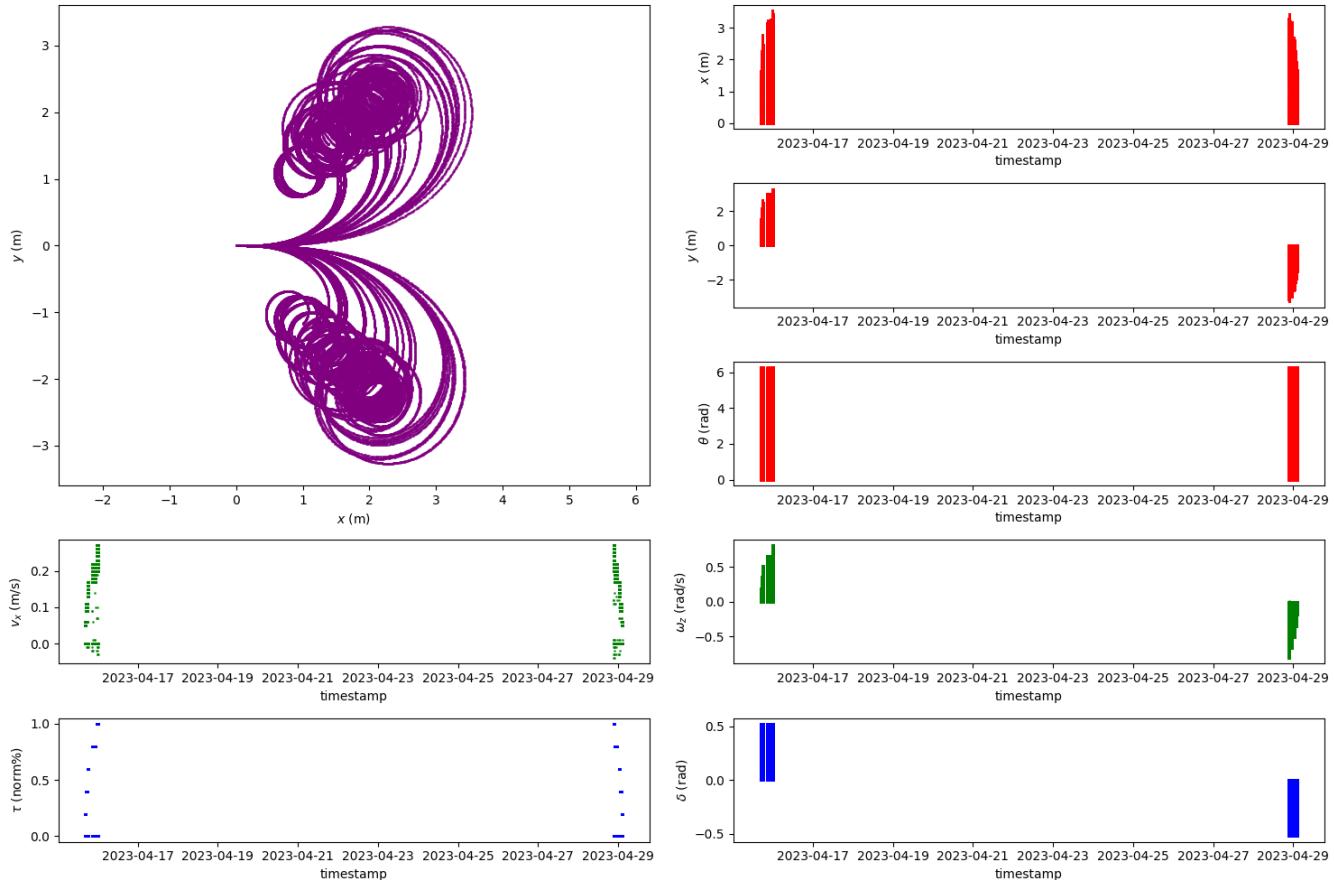
Model Inference - Skidpad Maneuver



Note: This work adheres to Markov assumption.

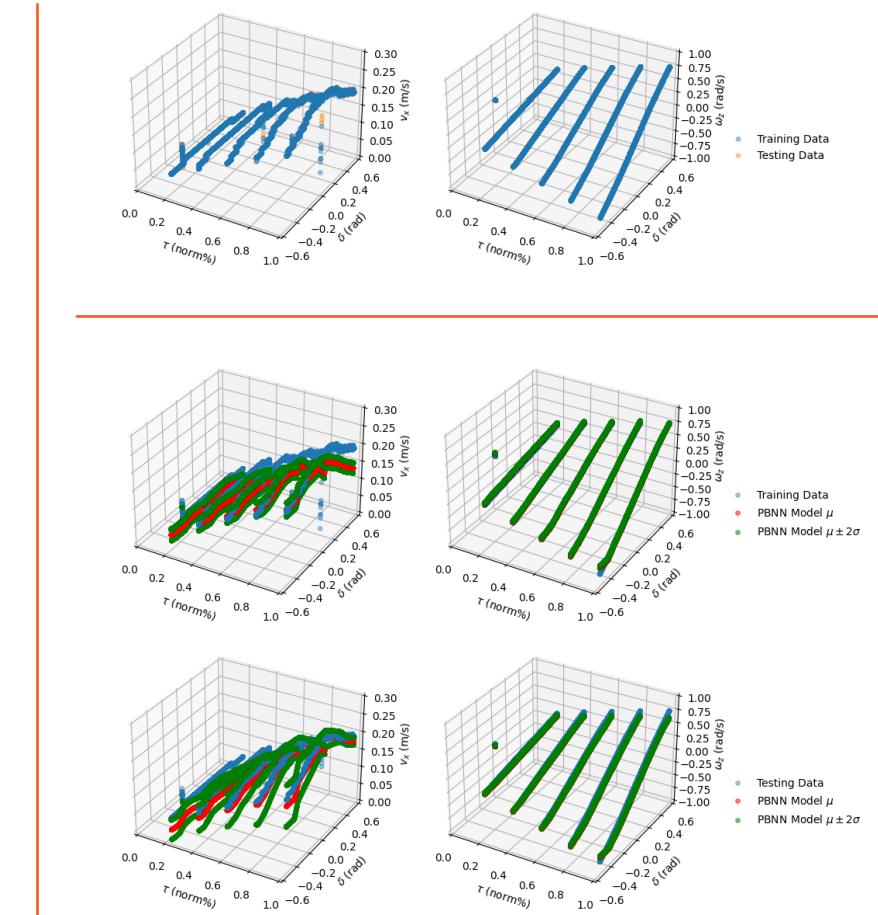
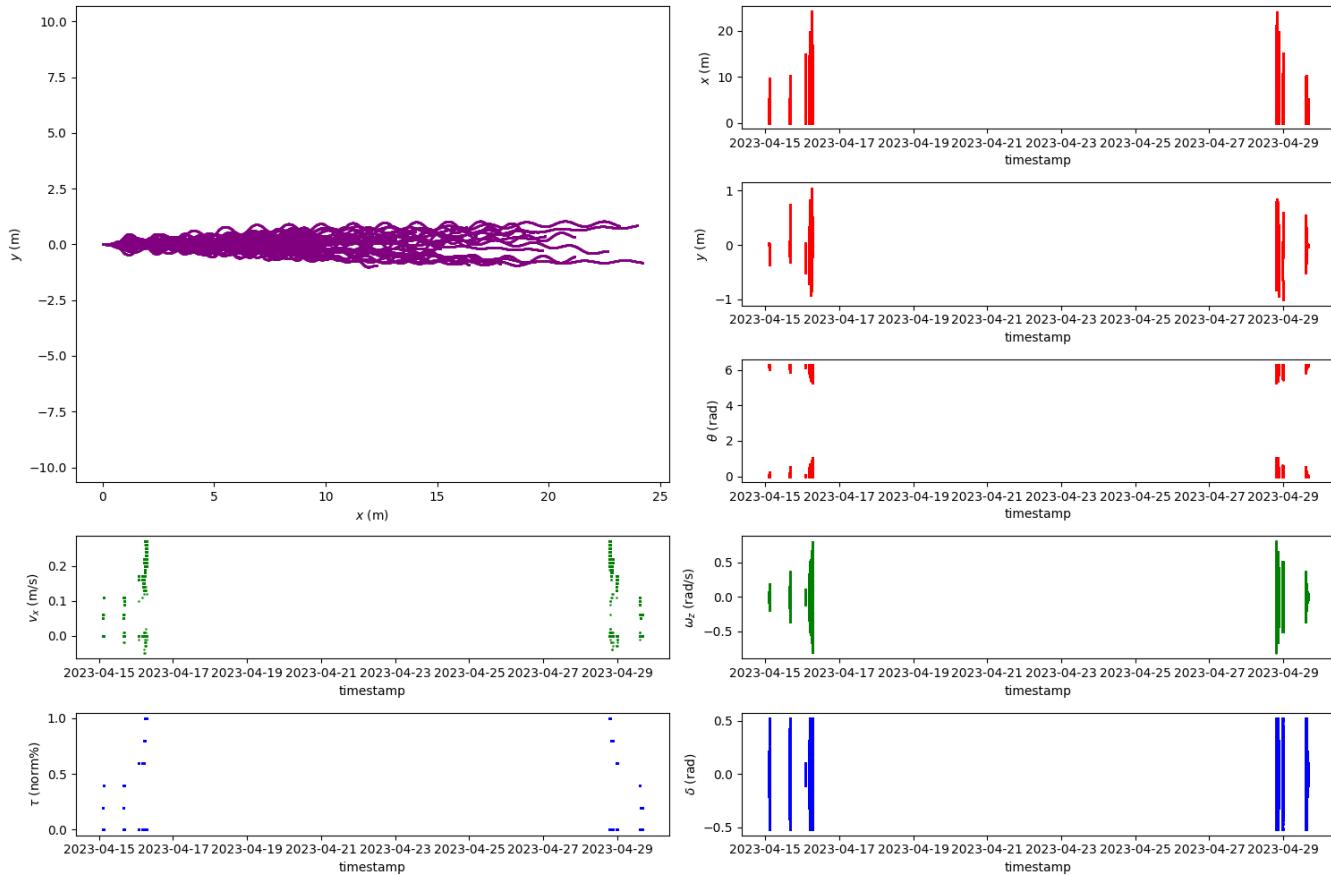


Model Inference - Fishhook Maneuver



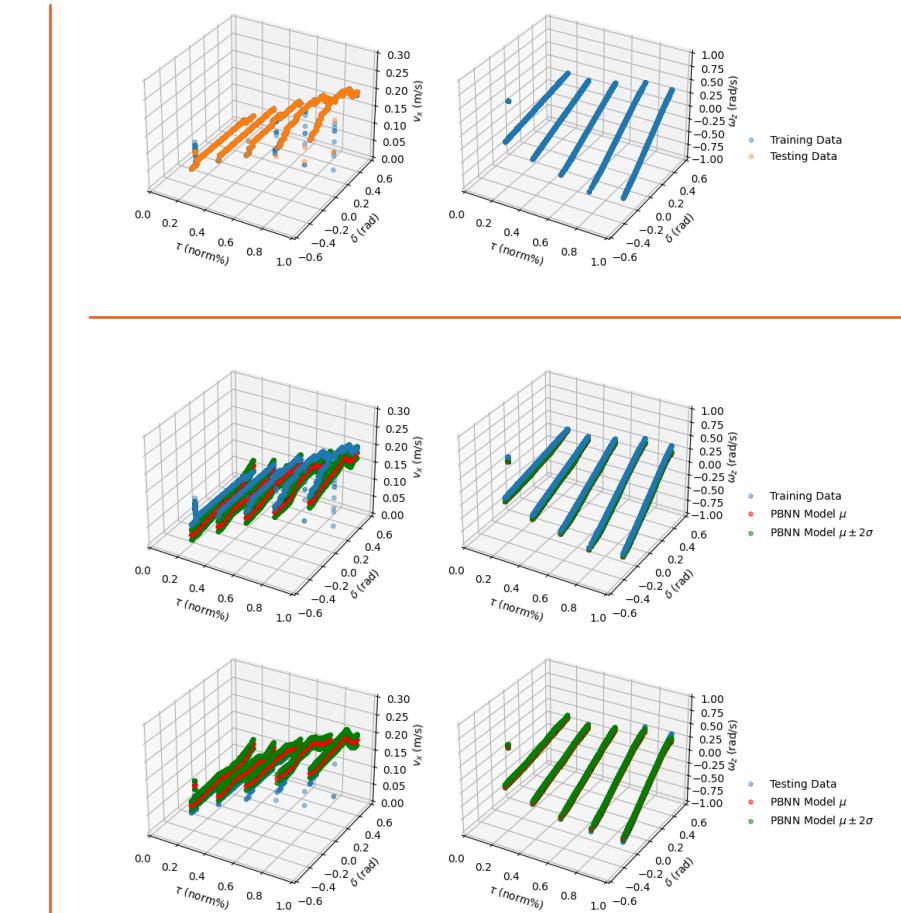
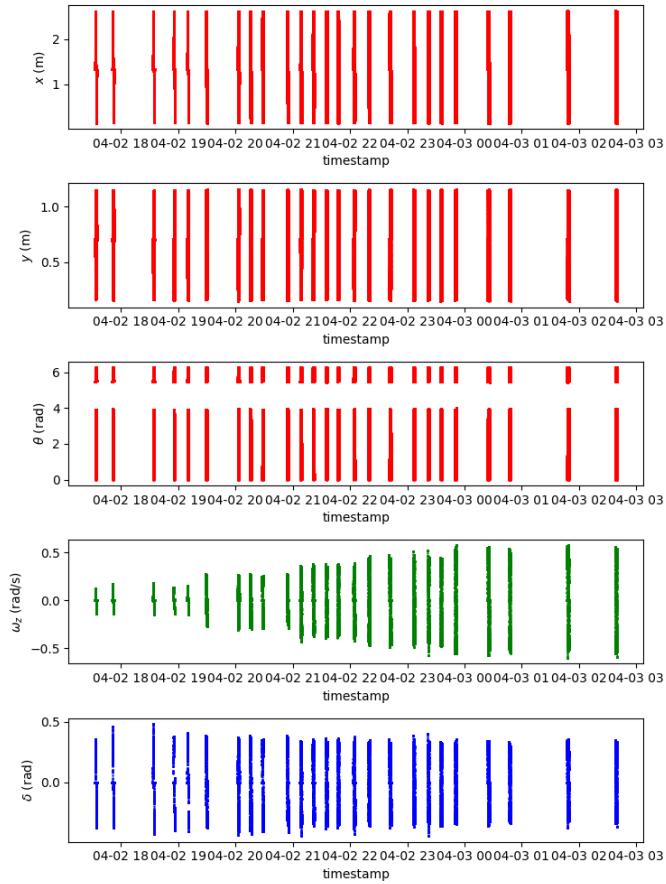
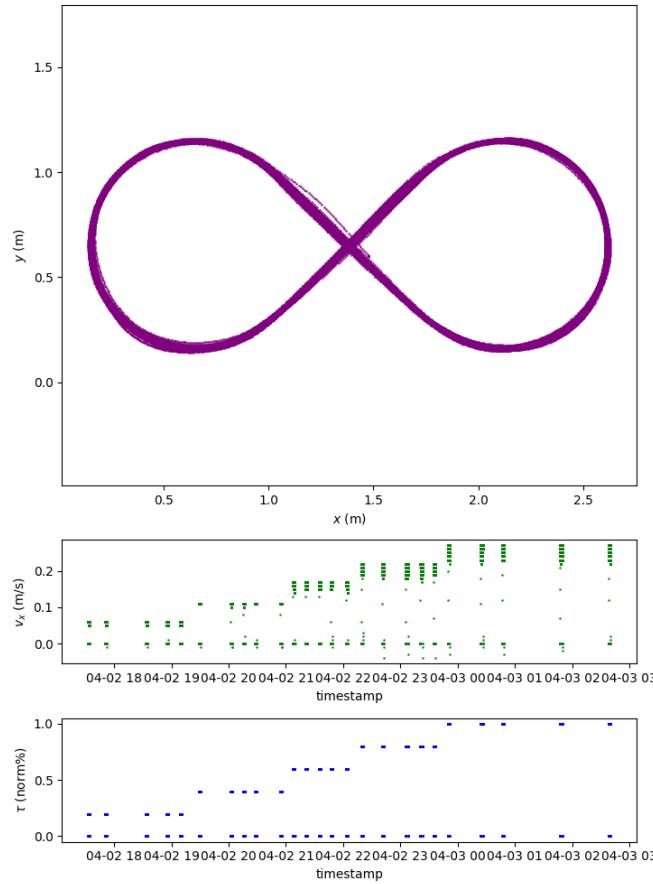
Note: This work adheres to Markov assumption.

Model Inference - Slalom Maneuver



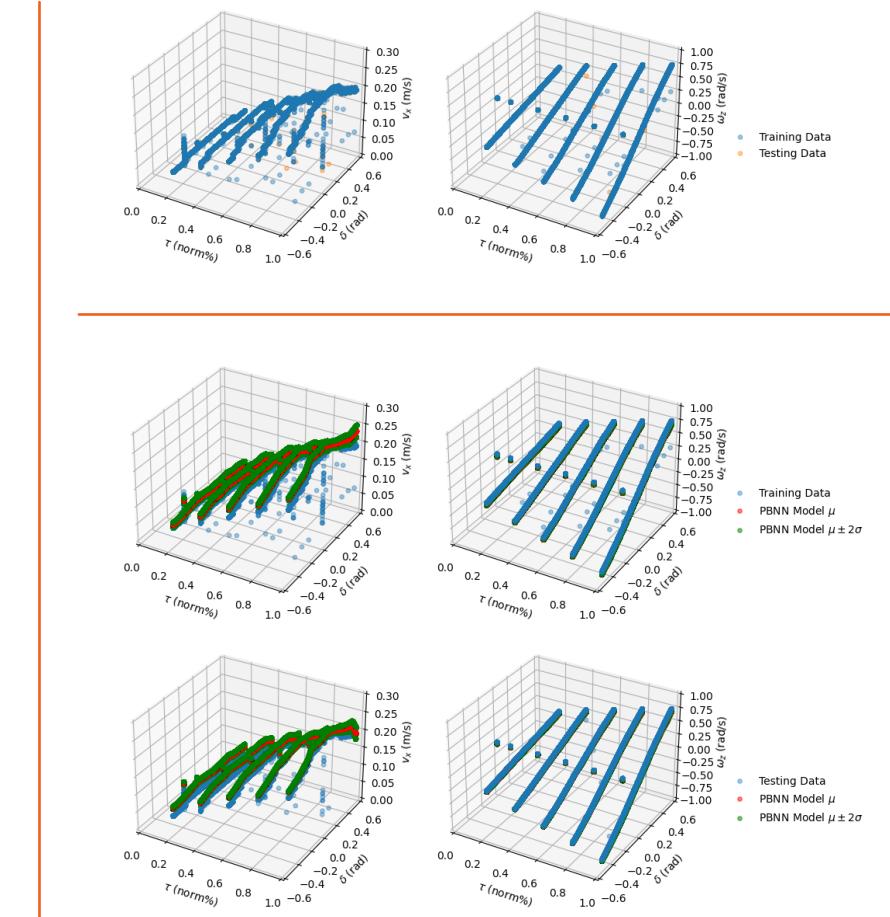
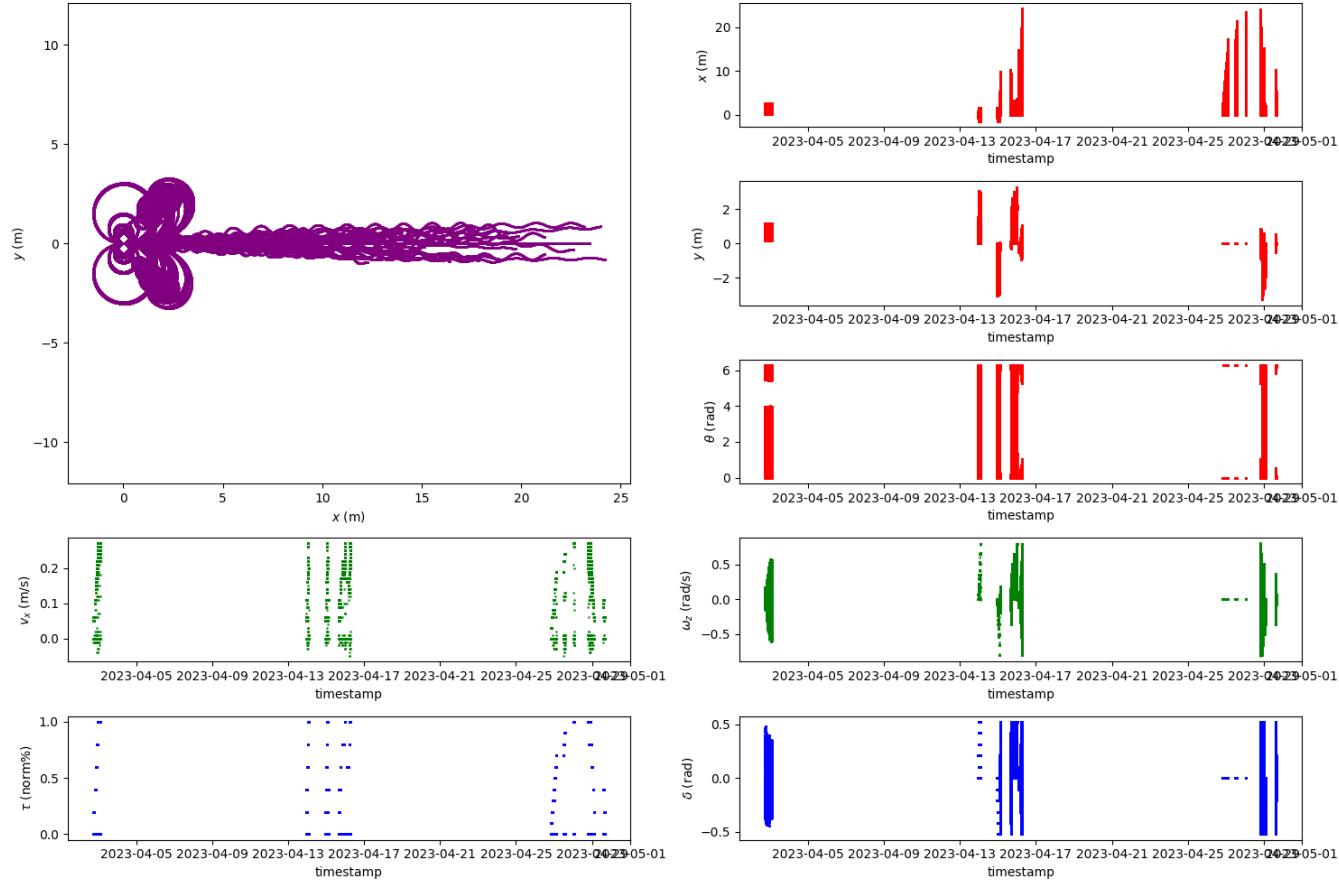
Note: This work adheres to Markov assumption.

Model Inference - Eight Maneuver



Note: This work adheres to Markov assumption.

Model Inference - Generalization Across All Maneuvers



Note: This work adheres to Markov assumption.



Future Plans

- Formulate and implement probabilistic Bayesian LSTM (delayed embedding)
- Try different input representation technique (e.g. timeseries snapshot of state information)
- Utilize other state (and perception) information from dataset
- Transfer learning for expanding model's ODD to different operating points
- Ensemble models for different operating points
- Simulation + real-world vehicle data (across scales)
- Active sampling to collect (sample) new data points smartly based on epistemic uncertainty



Thank You!

Open to questions and suggestions...

